CLOUDFLARE

Speedrun through Splicing Sockets with SOCKMAP

Jakub Sitnicki
Systems Engineer
Cloudflare

Cilium + eBPF Day EUROPE

19 March 2024 | Paris, France

$ whoami

Linux / OS Team @ Cloudflare

🌽 roll out fresh kernels

🐞 squash bugs

🕵️ troubleshoot stuff

✨ prototype features

$ whoami

Linux / OS Team @ Cloudflare

🌽 roll out fresh kernels

🐞 squash bugs

🕵️ troubleshoot stuff

✨ prototype features

SOCKMAP co-maintainer @ Linux upstream

🤏 small-time (= feature) maintainer

🪲 fix bugs

🧐 review patches

🤔 answer questions

About this talk

CLOUDFLARE

Good to know:

- ❏ network programming (`socket`, `connect`, `sendmsg`, `recvmsg`)
- ❏ basics of eBPF (what are BPF maps, programs, hooks, `bpftool`)
- ❏ building blocks of containers (cgroups, namespaces)

Goals:

- ❏ know that SOCKMAP exists
- ❏ have idea how / when / what for use it
- ❏ feel ready to dive deeper

CLOUDFLARE

Agenda

1. What can SOCKMAP do?
2. What is SOCKMAP?
3. How to set up SOCKMAP?
4. How to get sockets into a SOCKMAP?
5. Supported socket splicing setups
6. Real life use cases

Section

What can SOCKMAP do for you?

What can SOCKMAP do for… container networking

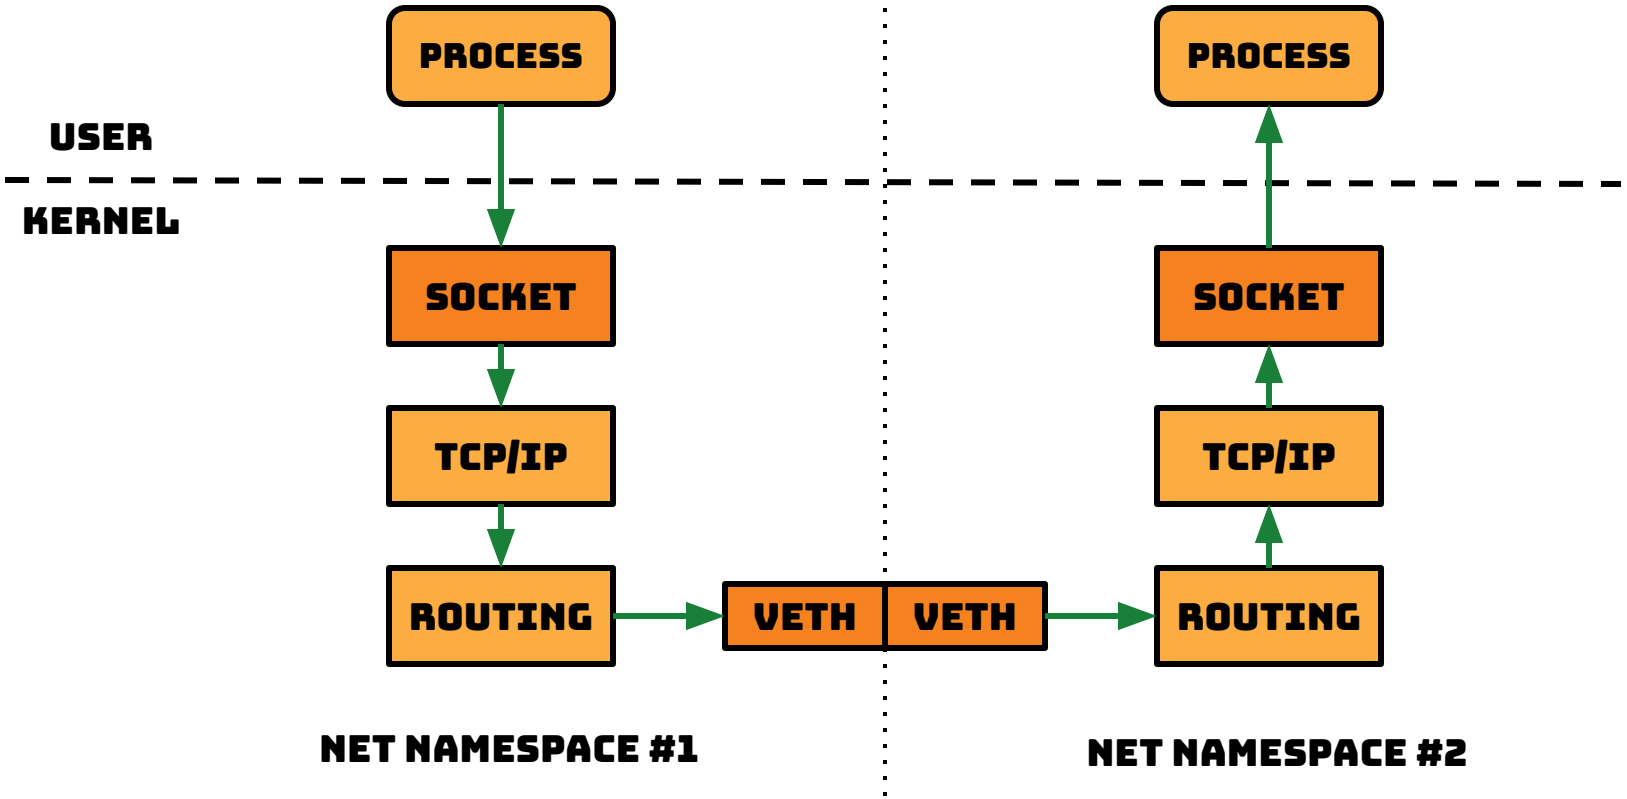

What can SOCKMAP do for… container networking

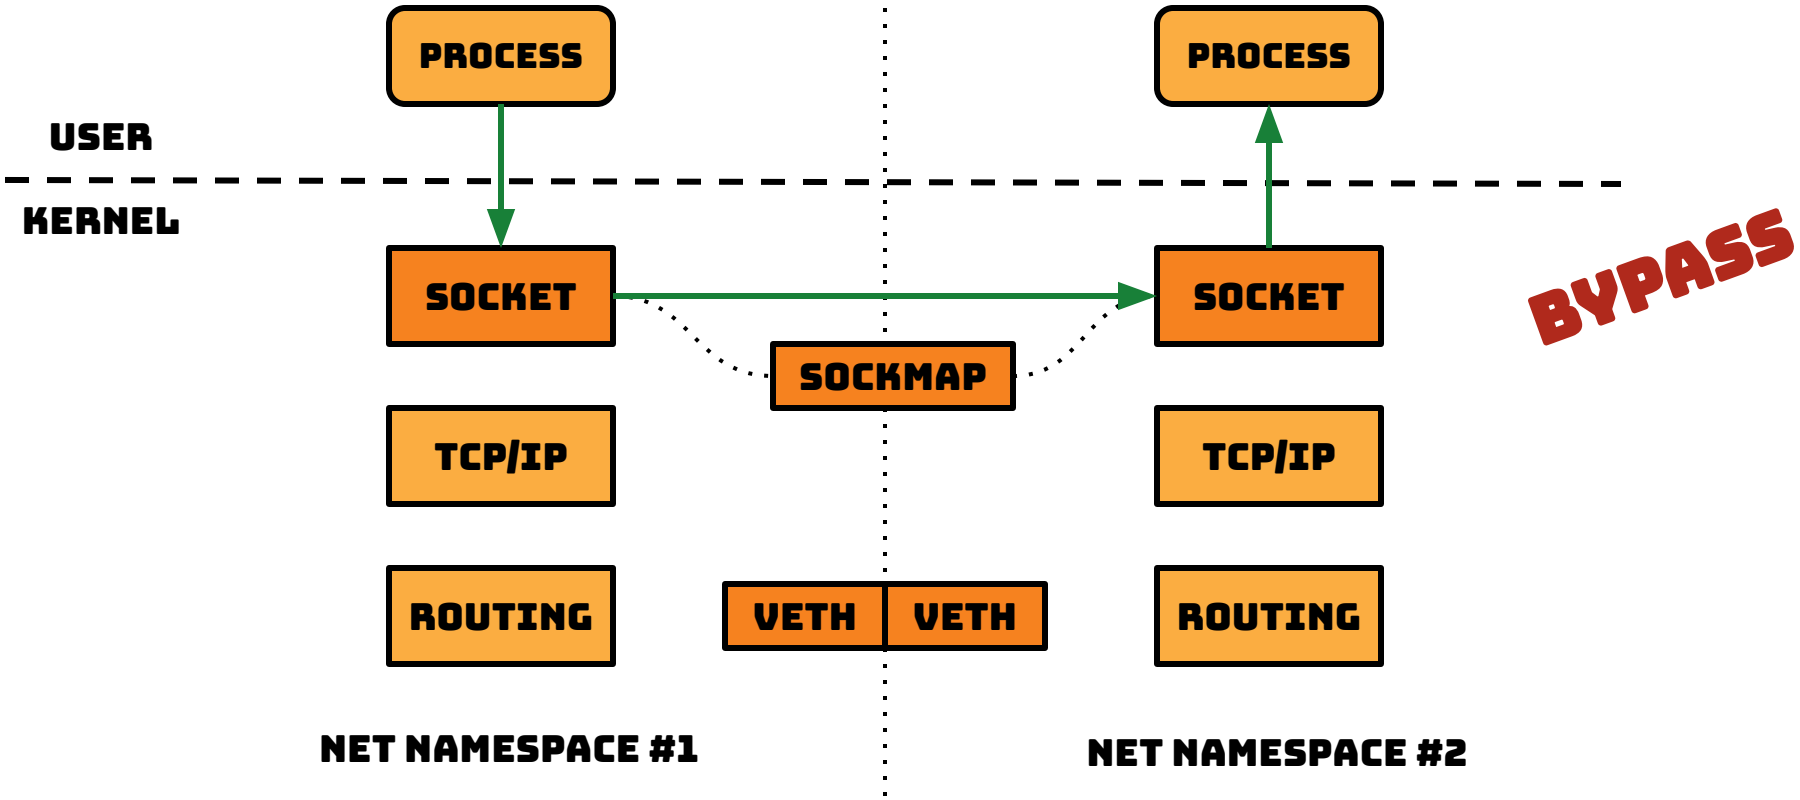

Let’s set it up!

Create two network namespaces

```
# ip netns add A
# ip netns add B
```

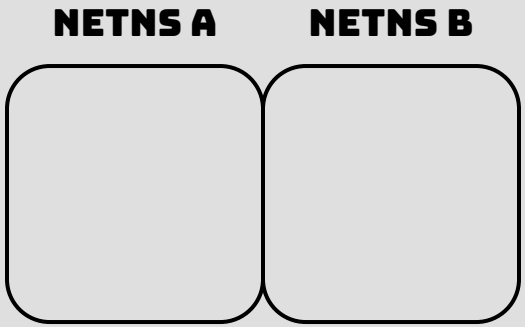

Let's set it up!

Create two network namespaces

```
# ip netns add A
# ip netns add B
```

Link network namespaces with a veth pair

```
# ip -n A link add name veth1 type veth \
                peer name veth2 netns B
```

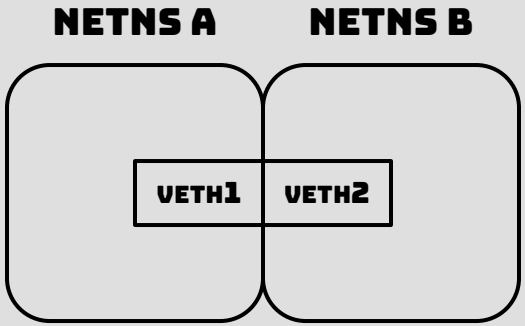

Let's set it up!

CLOUDFLARE

Create two network namespaces

```
# ip netns add A
# ip netns add B
```

Link network namespaces with a veth pair

```
# ip -n A link add name veth1 type veth \
                peer name veth2 netns B
```

Bring up the links inside network namespaces

```
# ip -n A link set dev veth1 up
# ip -n B link set dev veth2 up
```

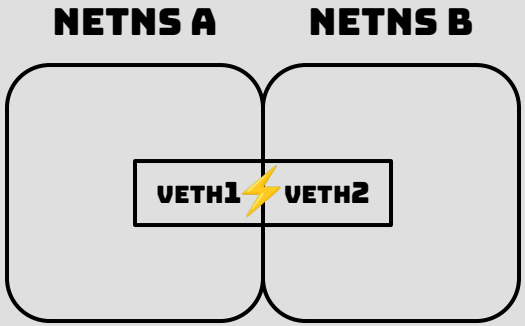

Let's set it up!

CLOUDFLARE

Create two network namespaces

```
# ip netns add A
# ip netns add B
```

Link network namespaces with a veth pair

```
# ip -n A link add name veth1 type veth \
                peer name veth2 netns B
```

Bring up the links inside network namespaces

```
# ip -n A link set dev veth1 up
# ip -n B link set dev veth2 up
```

Assign addresses to links inside network namespaces

```
# ip -n A addr add 10.0.0.1/24 dev veth1
# ip -n B addr add 10.0.0.2/24 dev veth2
```

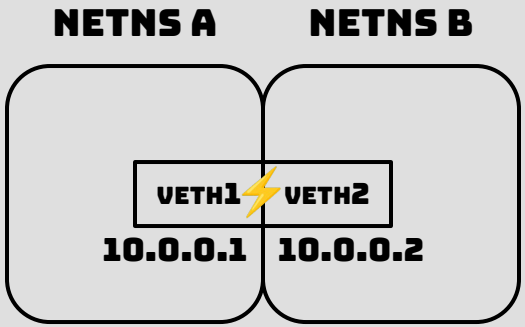

Measure latency, no SOCKMAP first

Run TCP server in netns A

```
# ip netns exec A \
  sockperf server -i 10.0.0.1 --tcp --daemonize
```

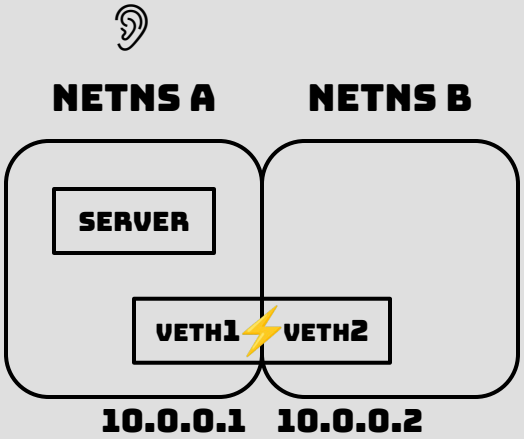

Measure latency, no SOCKMAP first

Run TCP server in netns A

```
# ip netns exec A \
  sockperf server -i 10.0.0.1 --tcp --daemonize
```

Run TCP client in netns B

```
# ip netns exec B \
  sockperf ping-pong -i 10.0.0.1 --tcp --time 30
```

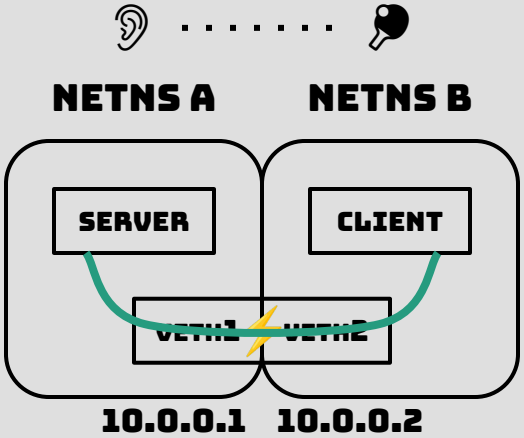

Measure latency, no SOCKMAP first

```
Run TCP server in netns A
# ip netns exec A \
  sockperf server -i 10.0.0.1 --tcp --daemonize


Run TCP client in netns B
# ip netns exec B \
  sockperf ping-pong -i 10.0.0.1 --tcp --time 30
…
sockperf: [Total Run] RunTime=30.000 sec; Warm up time=400 msec; SentMessages=2599753;
ReceivedMessages=2599752
…
sockperf: ====> avg-latency=5.748 (std-dev=2.010, mean-ad=0.322, median-ad=0.220,
siqr=0.239, cv=0.350, std-error=0.001, 99.0% ci=[5.745, 5.751])
sockperf: # dropped messages = 0; # duplicated messages = 0; # out-of-order messages = 0
sockperf: Summary: Latency is 5.748 usec
```

5.8 ± 2.0 μsec

Set up SOCKMAP bypass

Load BPF programs and create BPF maps

```
# bpftool prog loadall redir_bypass.bpf.o /sys/fs/bpf \
                pinmaps /sys/fs/bpf
```

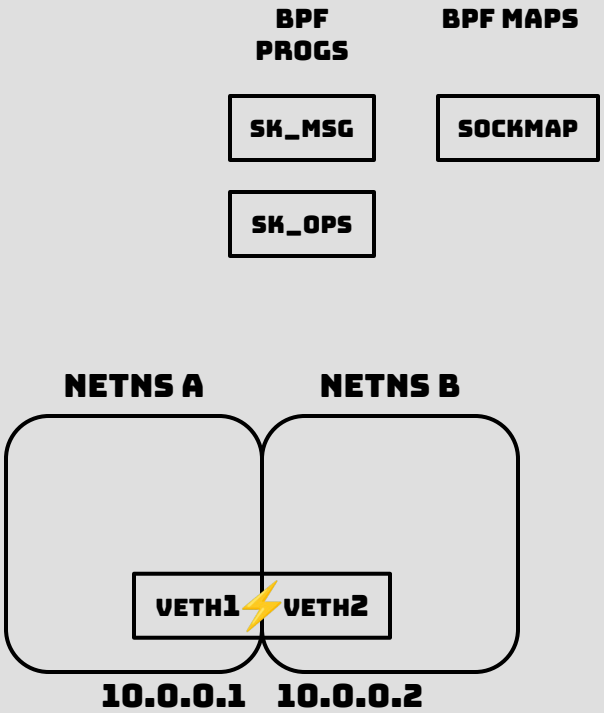

Set up SOCKMAP bypass

CLOUDFLARE

Load BPF programs and create BPF maps

```
# bpftool prog loadall redir_bypass.bpf.o /sys/fs/bpf \
                pinmaps /sys/fs/bpf
```

Attach BPF program to BPF map

```
# bpftool prog attach \
  pinned /sys/fs/bpf/sk_msg_prog sk_msg_verdict \
  pinned /sys/fs/bpf/sock_map
```

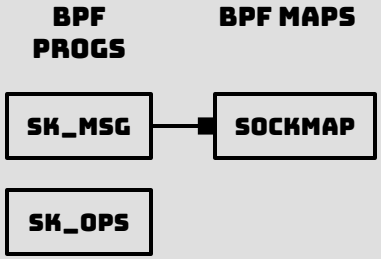

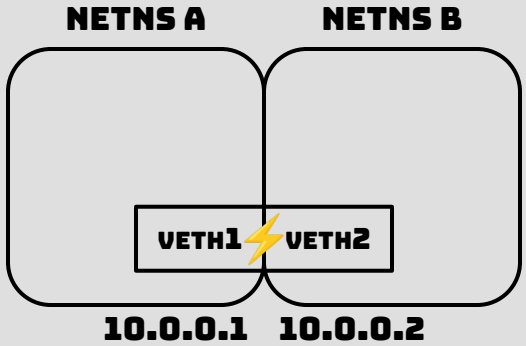

Set up SOCKMAP bypass

Load BPF programs and create BPF maps

```
# bpftool prog loadall redir_bypass.bpf.o /sys/fs/bpf \
                pinmaps /sys/fs/bpf
```

Attach BPF program to BPF map

```
# bpftool prog attach \
  pinned /sys/fs/bpf/sk_msg_prog sk_msg_verdict \
  pinned /sys/fs/bpf/sock_map
```

Create a test cgroup

```
# mkdir /sys/fs/cgroup/unified/test.slice
```

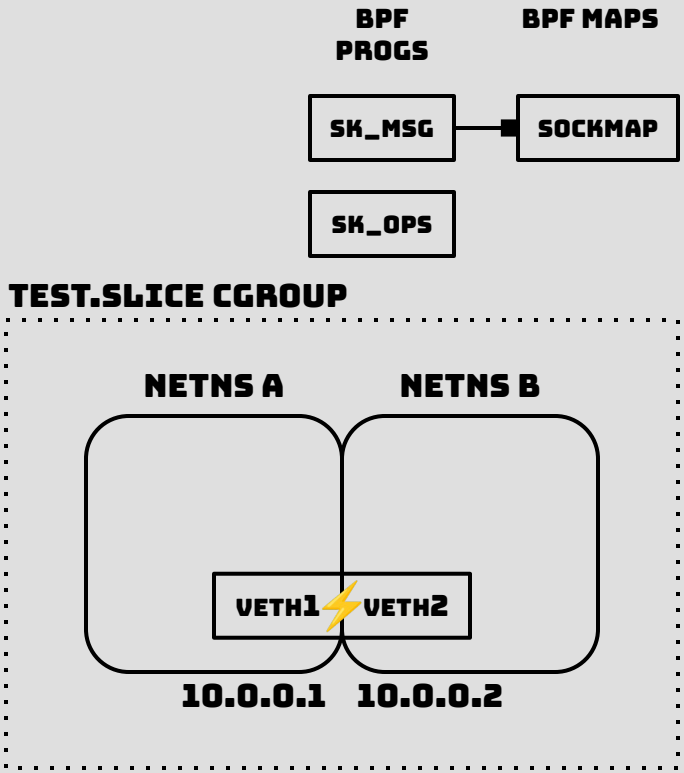

Set up SOCKMAP bypass

CLOUDFLARE

Load BPF programs and create BPF maps

```
# bpftool prog loadall redir_bypass.bpf.o /sys/fs/bpf \
                pinmaps /sys/fs/bpf
```

Attach BPF program to BPF map

```
# bpftool prog attach \
  pinned /sys/fs/bpf/sk_msg_prog sk_msg_verdict \
  pinned /sys/fs/bpf/sock_map
```

Create a test cgroup

```
# mkdir /sys/fs/cgroup/unified/test.slice
```

Attach BPF program to cgroup

```
# bpftool cgroup attach \
  /sys/fs/cgroup/unified/test.slice \
  cgroup_sock_ops pinned /sys/fs/bpf/sockops_prog
```

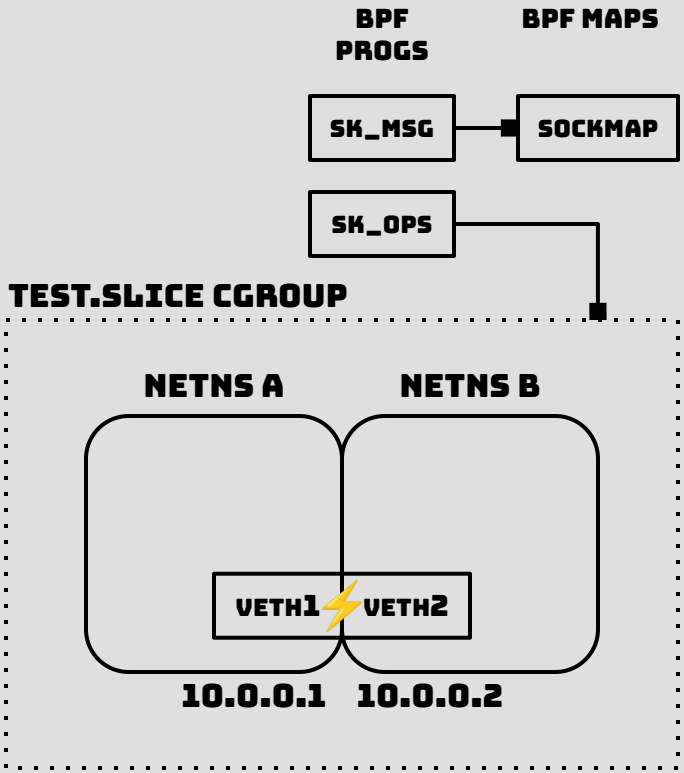

Repeat test with SOCKMAP bypass

Spawn client and server inside the test cgroup

```
# echo $$ > /sys/fs/cgroup/unified/test.slice/cgroup.procs
```

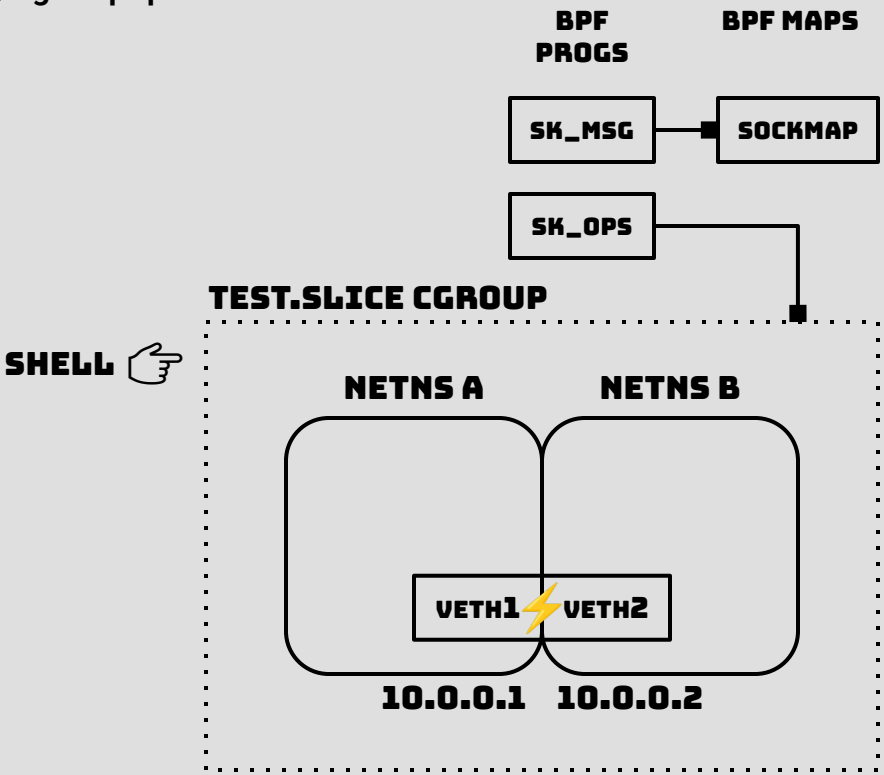

Repeat test with SOCKMAP bypass

Spawn client and server inside the test cgroup

```
# echo $$ > /sys/fs/cgroup/unified/test.slice/cgroup.procs
```

Run TCP server in netns A

```
# ip netns exec A \
  sockperf server -i 10.0.0.1 --tcp --daemonize
```

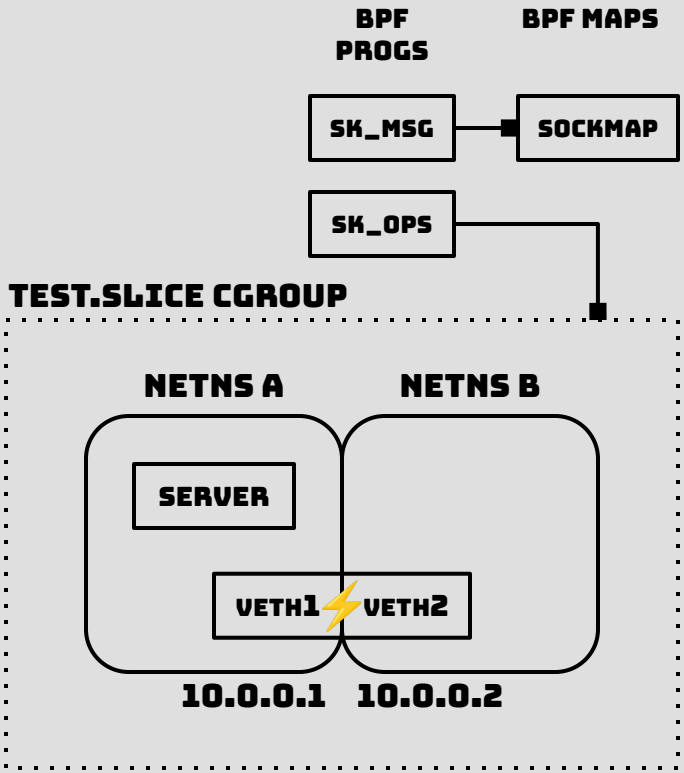

Repeat test with SOCKMAP bypass

Spawn client and server inside the test cgroup

```
# echo $$ > /sys/fs/cgroup/unified/test.slice/cgroup.procs
```

Run TCP server in netns A

```
# ip netns exec A \
  sockperf server -i 10.0.0.1 --tcp --daemonize
```

Run TCP client

```
# ip netns exec B \
   sockperf ping-pong -i 10.0.0.1 --tcp --time 30
```

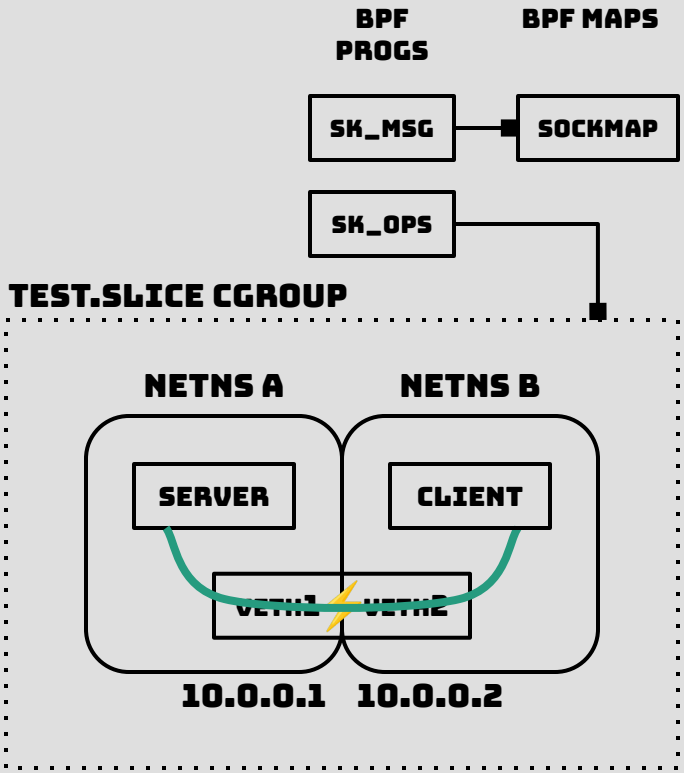

Repeat test with SOCKMAP bypass

CLOUDFLARE

```
Spawn client and server inside the test cgroup
# echo $$ > /sys/fs/cgroup/unified/test.slice/cgroup.procs

Run TCP server in netns A
# ip netns exec A \
  sockperf server -i 10.0.0.1 --tcp --daemonize

Run TCP client
# ip netns exec B \
   sockperf ping-pong -i 10.0.0.1 --tcp --time 30
sockperf: [Total Run] RunTime=30.000 sec; Warm up time=400 msec; SentMessages=3189584;
ReceivedMessages=3189583
…
sockperf: ====> avg-latency=4.686 (std-dev=2.862, mean-ad=0.250, median-ad=0.216,
siqr=0.173, cv=0.611, std-error=0.002, 99.0% ci=[4.682, 4.690])
sockperf: # dropped messages = 0; # duplicated messages = 0; # out-of-order messages = 0
sockperf: Summary: Latency is 4.686 usec
```

4.7 ± 2.9 μsec

Without and with SOCKMAP bypass

CLOUDFLARE

before: 5.8 ± 2.0 μsec

↓ - 18.5%

after: 4.7 ± 2.9 μsec

Run the benchmark yourself:
https://github.com/jsitnicki/kubecon-2024-sockmap/tree/main/examples/redir-bypass

Section

What is SOCKMAP?

What is SOCKMAP?

Two things

What is SOCKMAP?

CLOUDFLARE

Collection / container for socket references in Linux kernel

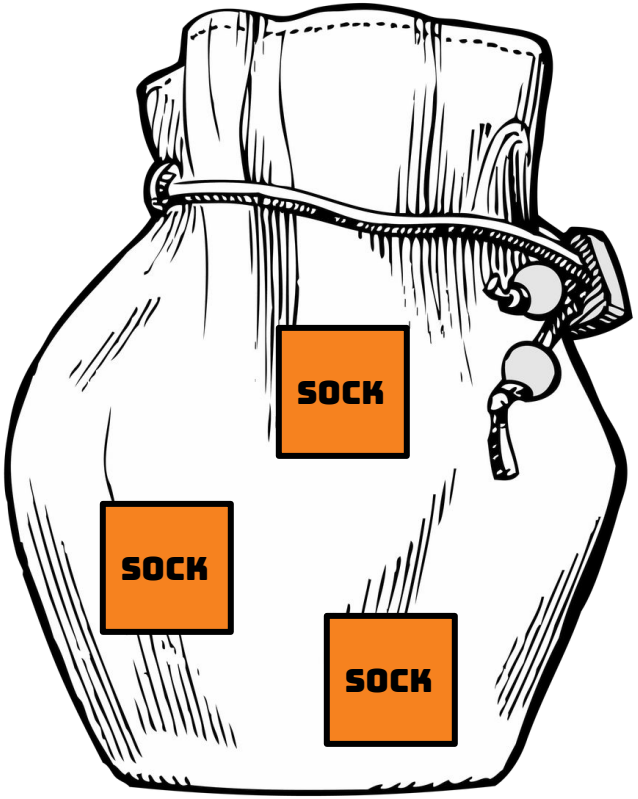

What is SOCKMAP?

CLOUDFLARE

SOCKMAP API

1. **container for sockets**

- BPF map (K/V store)
- holds weak refs to sockets

What is SOCKMAP?

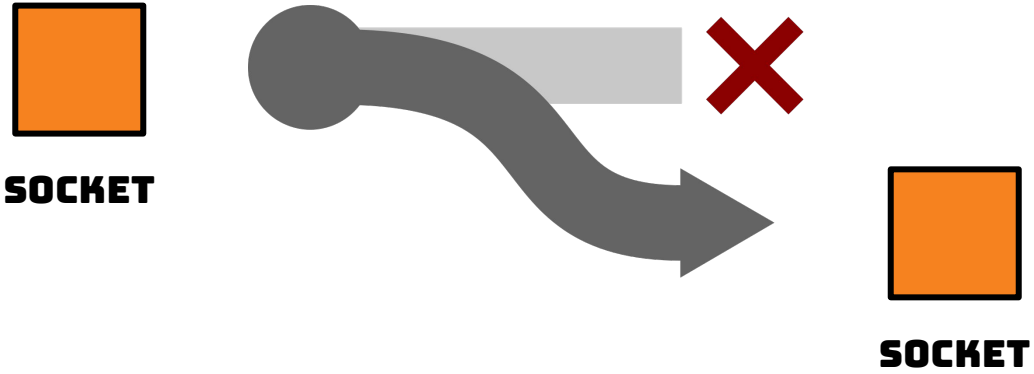

API for enforcing policy
and
redirecting data between sockets

What is SOCKMAP?

CLOUDFLARE

SOCKMAP API

1. **container for sockets**

- BPF map (K/V store)
- holds weak refs to sockets

2. **policy enforcement & redirecting packets**

- BPF programs to filter or redirect (steer) data from socket to socket
- hooks into socket layer

Section

How to set up SOCKMAP?

① Open a connected (established) socket

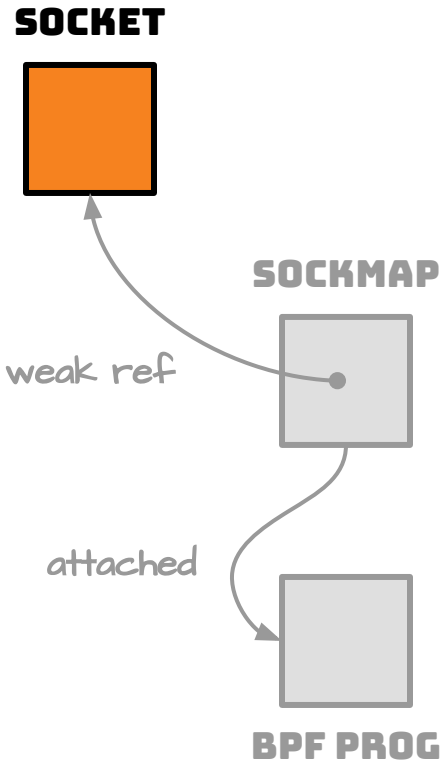

active open

```
socket(AF_INET, SOCK_STREAM, IPPROTO_IP) = 8
connect(8, {sa_family=AF_INET,
            sin_port=htons(41895),
            sin_addr=inet_addr("127.0.0.1")}, 16) = 0
```

passive open

```
socket(AF_INET, SOCK_STREAM, IPPROTO_IP) = 7
bind(7, {sa_family=AF_INET,
         sin_port=htons(41895),
         sin_addr=inet_addr("127.0.0.1")}, 16) = 0
listen(7, 4096)                           = 0
accept(7, NULL, NULL)                     = 9
```

What sockets can you use?

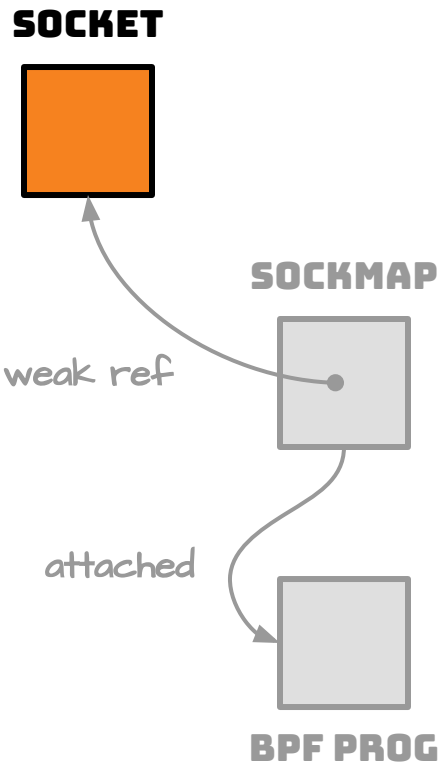

connected (established) socket:

- ❑ **TCP**
- ❑ **UDP**
- ❑ **UNIX** (STREAM, DGRAM)
- ❑ **VSOCK** (STREAM, SEQPACKET)

`man 7 {tcp,udp,unix,vsock}`

② Create a BPF map - SOCKMAP or SOCKHASH

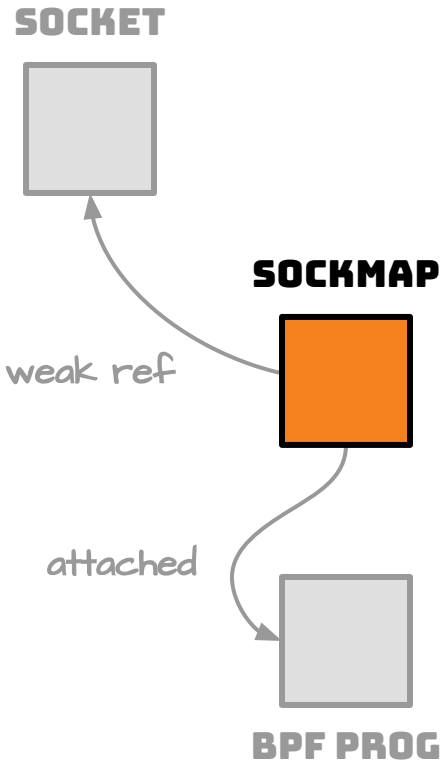

use the `bpf()` syscall

or a library wrapper (`ebpf-go`, `libbpf`)

```
bpf(BPF_MAP_CREATE, {map_type=BPF_MAP_TYPE_SOCKMAP,
                     key_size=4,
                     value_size=8,
                     max_entries=1,
                     map_flags=0,
                     …}, 72) = 5
```

② Create a BPF map

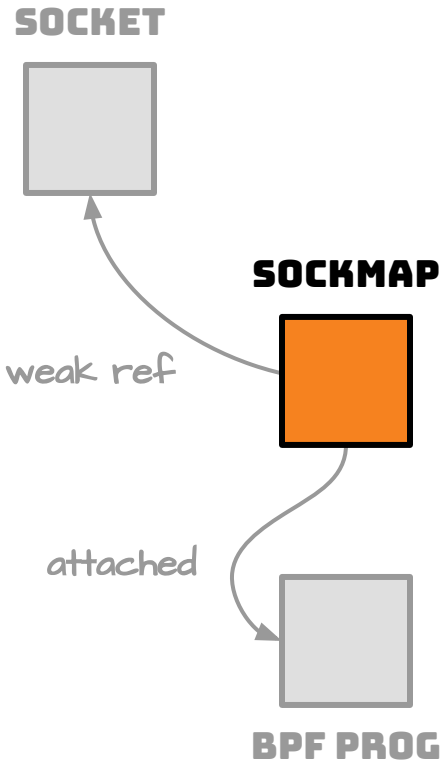

or use `bpftool map create` command

```
bpftool map create                                          \
  /sys/fs/bpf/sockmap `# path on bpffs`                     \
  type sockmap        `# sockmap or sockhash`               \
  key 4               `# always 4 bytes for sockmap`        \
  value 8             `# use 8 bytes for dump to work`      \
  entries 1                                                 \
  name sockmap

bpftool map show pinned /sys/fs/bpf/sockmap
3: sockmap  name sockmap  flags 0x0
        key 4B  value 8B  max_entries 1  memlock 328B
```

What BPF maps can you use?

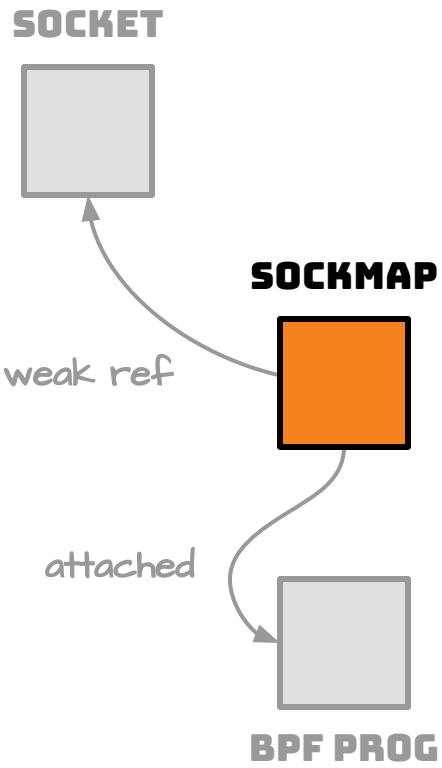

Map types:

- BPF_MAP_TYPE_**SOCKMAP**
  - **32-bit integer key**
- BPF_MAP_TYPE_**SOCKHASH**
  - **binary blob key**

Not to be confused with
BPF_MAP_TYPE_**REUSEPORT_SOCKARRAY**

https://docs.kernel.org/bpf/map_sockmap.html

③ Load a BPF program

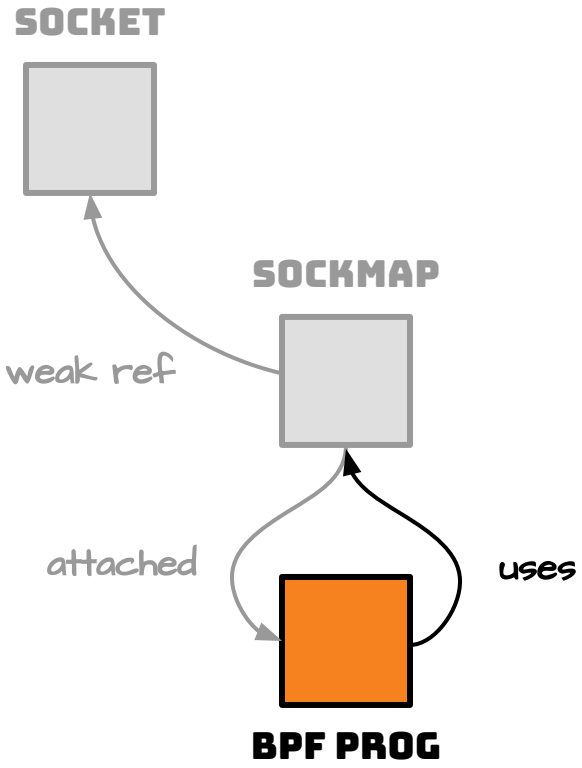


```
bpf(BPF_PROG_LOAD, {prog_type=BPF_PROG_TYPE_SK_MSG,
                    insn_cnt=6,
                    insns=0xcf2a70,
                    expected_attach_type=BPF_SK_MSG_VERDICT,
                    …}, 128) = 6
```

Program types:

- ❑ BPF_PROG_TYPE_**SK_MSG**
- ❑ BPF_PROG_TYPE_**SK_SKB**

③ Load a BPF program - it uses SOCKMAP

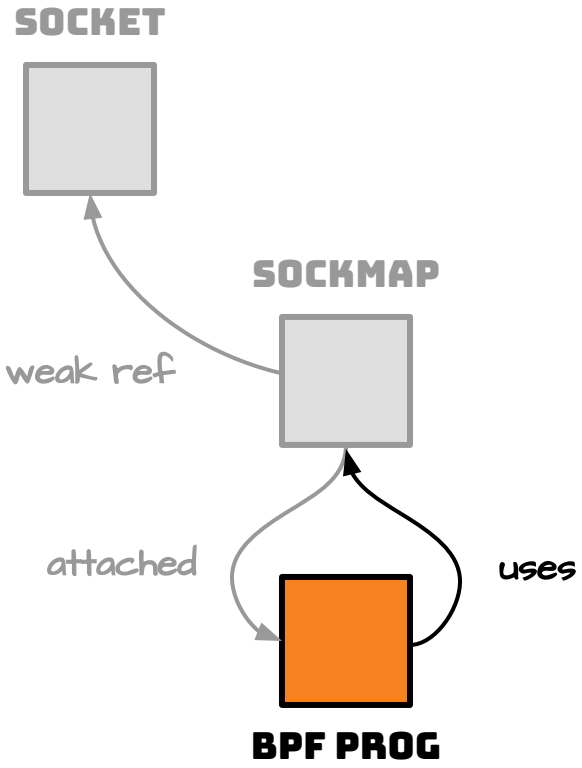


```
# bpftool prog dump xlated id 42
int prog_msg_redir_ingress(struct sk_msg_md * msg):
   0: (18) r2 = map[id:17]
   …
   5: (95) exit
# bpftool map show id 17
17: sockmap  name output  flags 0x0
        key 4B  value 8B  max_entries 1  memlock 328B
        pids sockmap-redir-m(331)
```

④ Attach BPF program to SOCKMAP

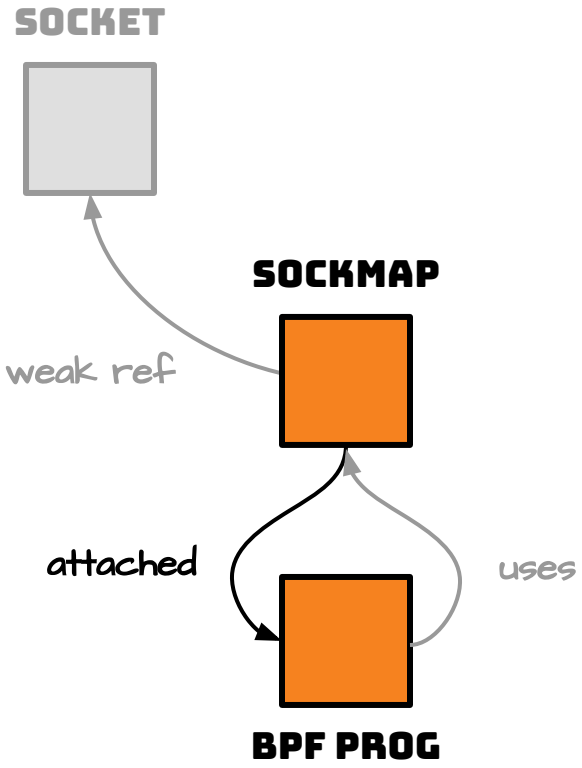


```
bpf(BPF_PROG_ATTACH, {target_fd=5,
                      attach_bpf_fd=6,
                      attach_type=BPF_SK_MSG_VERDICT,
                      attach_flags=0,
                      replace_bpf_fd=0}, 20) = 0
```

⑤ Insert socket into SOCKMAP

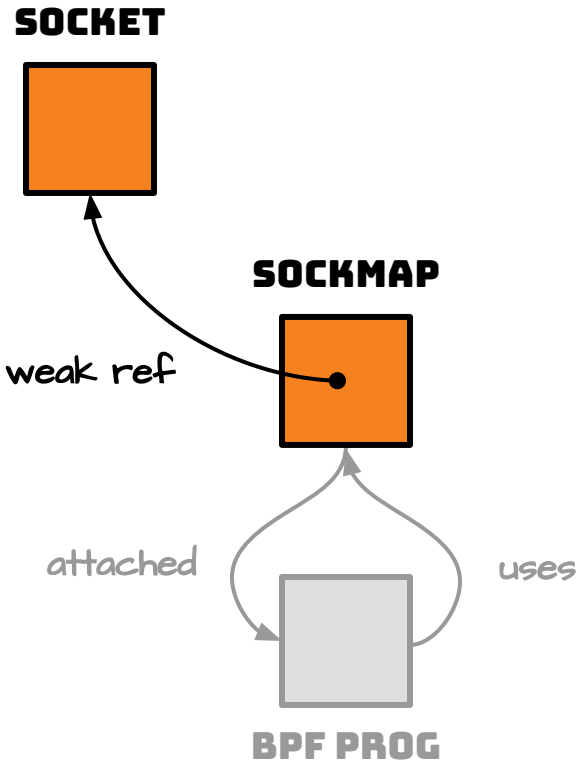


```
bpf(BPF_MAP_UPDATE_ELEM, {map_fd=5,
                          key=0x7ffeb3803870,
                          value=0x7ffeb3803868,
                          flags=BPF_NOEXIST}, 32) = 0
```

⚠ must be done after attaching the program

Section

How to get sockets into a SOCKMAP?

Easy case - Single process

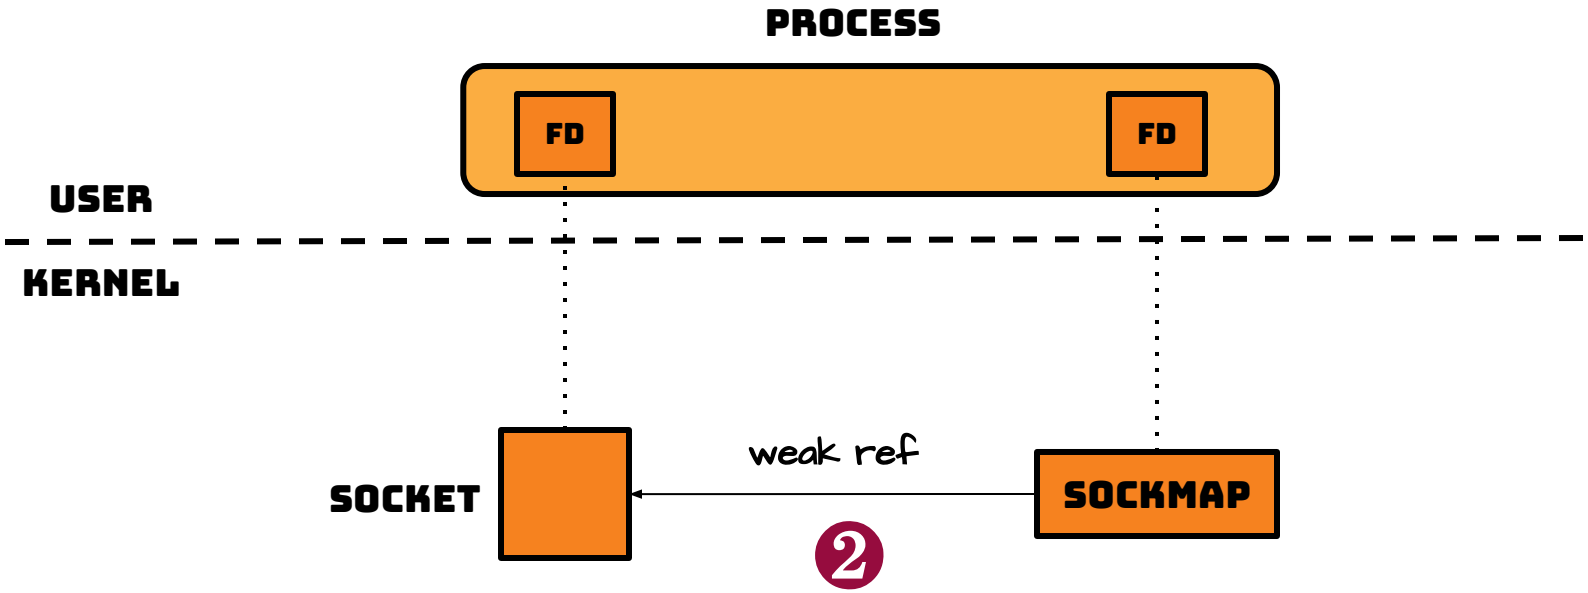

1 `bpf(BPF_MAP_UPDATE_ELEM, …)`

Socket FD handover with SCM_RIGHTS

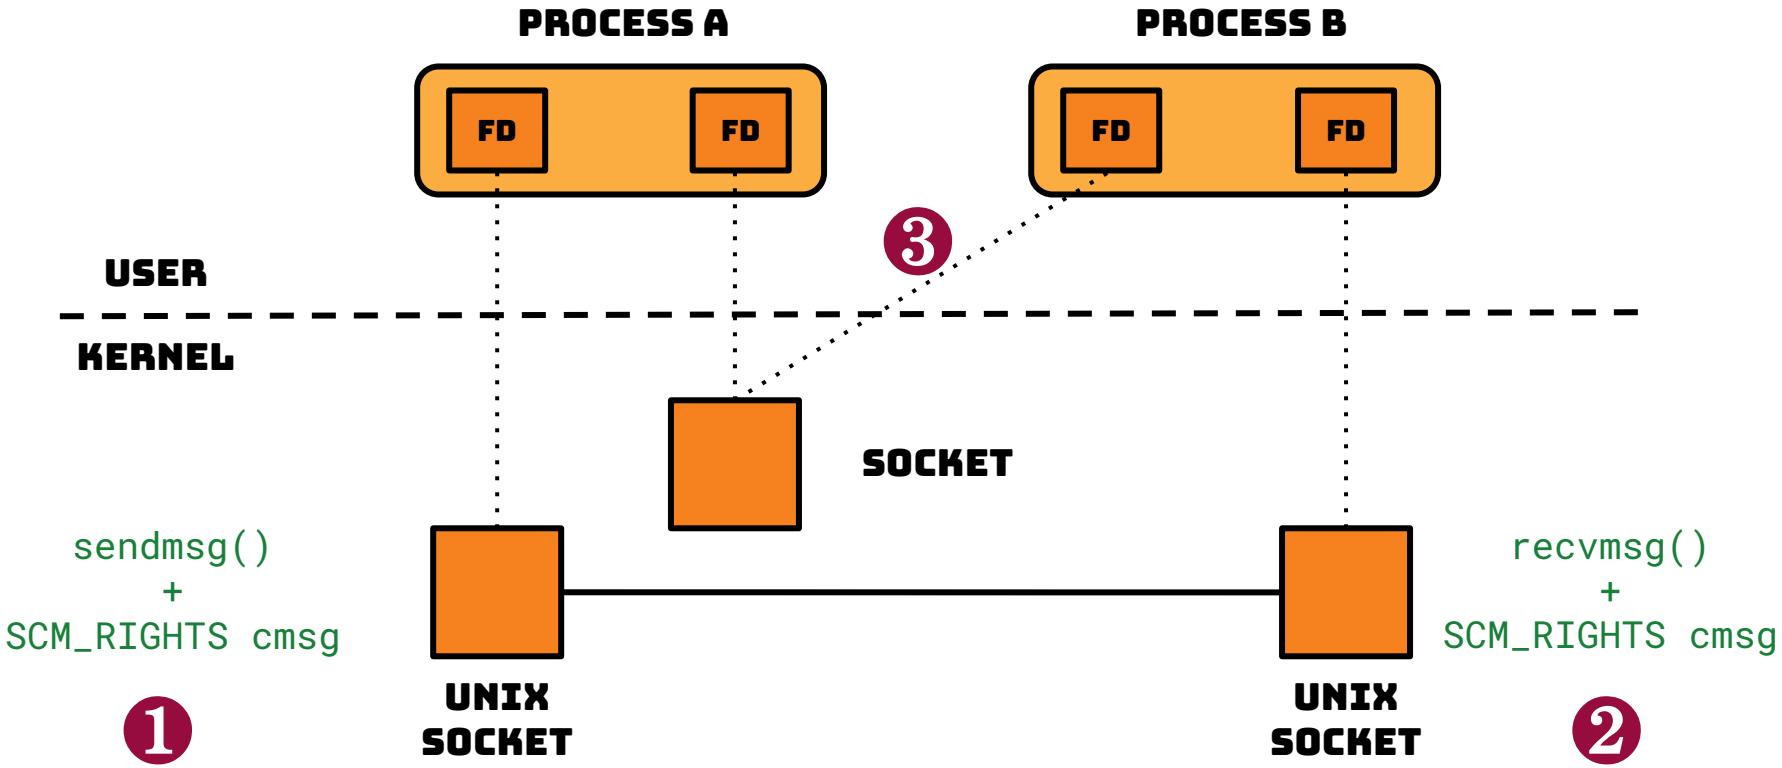

https://manpages.debian.org/unstable/manpages/unix.7.en.html#SCM_RIGHTS
https://blog.cloudflare.com/know-your-scm_rights/

“Steal” a socket FD

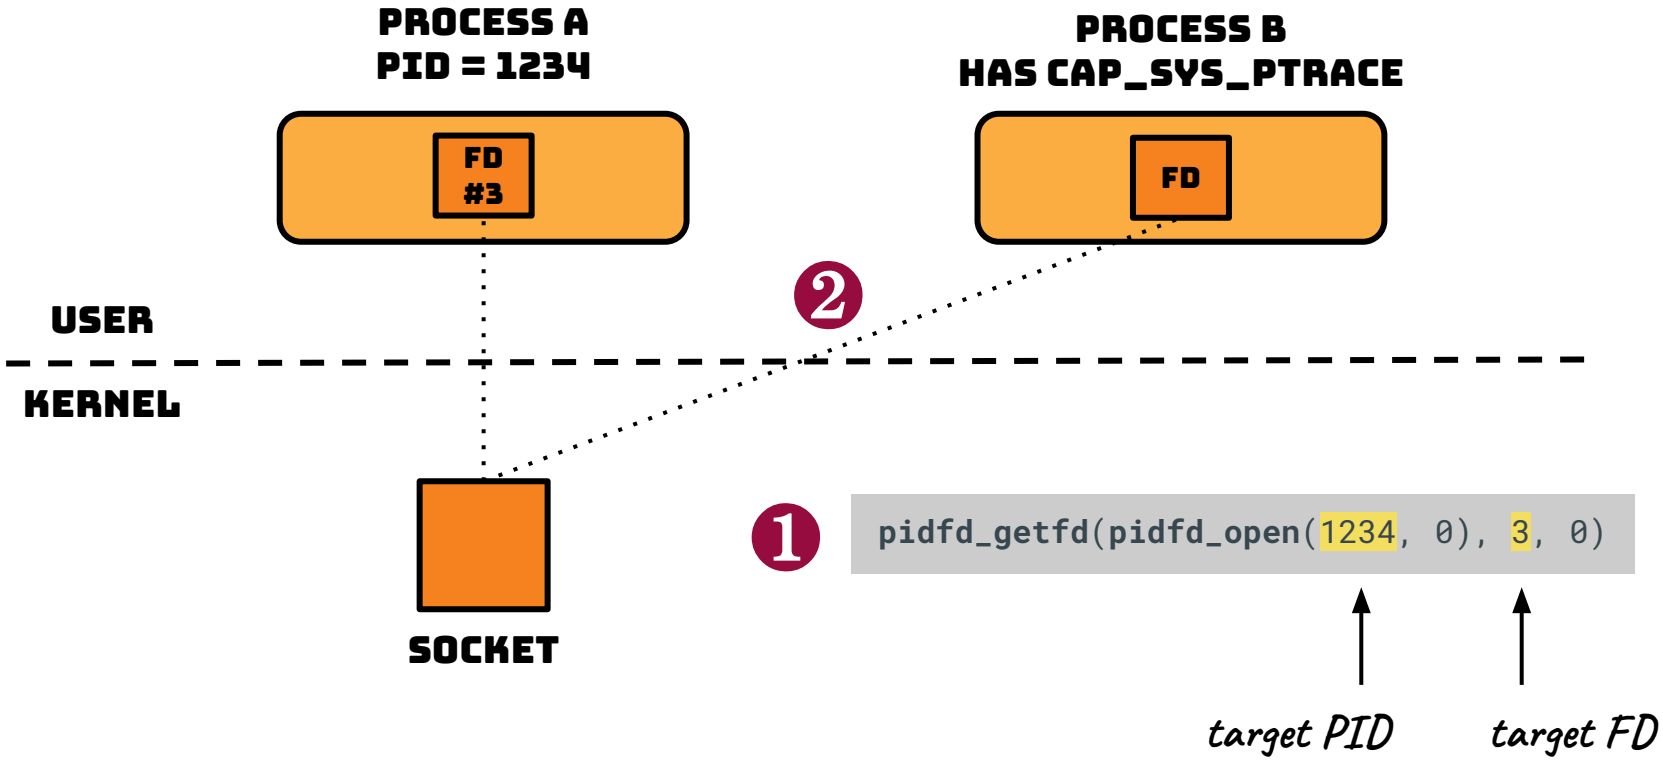

Linux 5.6+, requires CAP_SYS_PTRACE

https://manpages.debian.org/unstable/manpages-dev/pidfd_getfd.2.en.html

BPF sock_ops program attached to cgroup (TCP only)

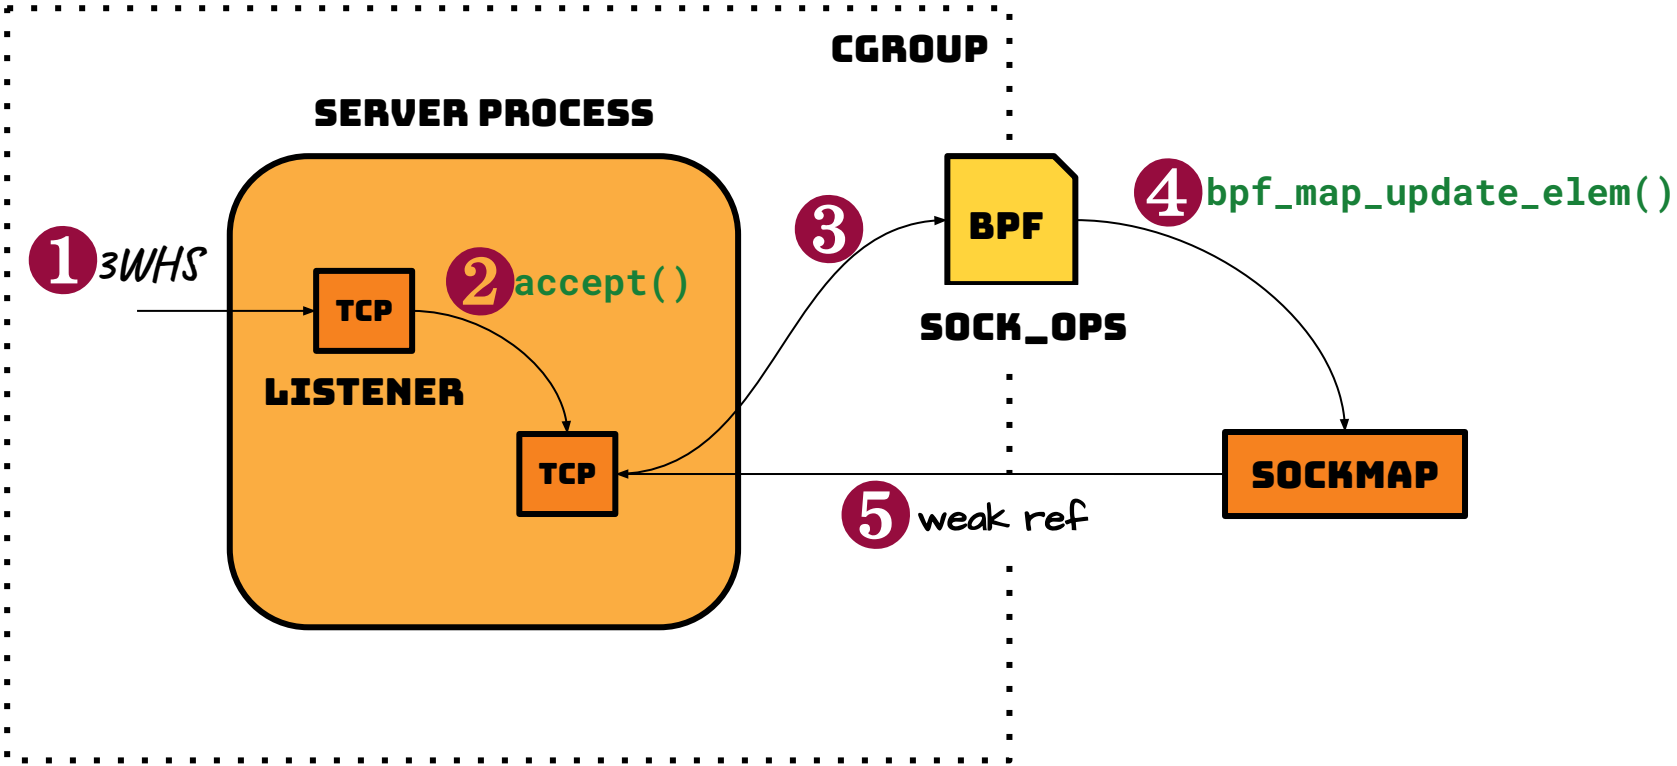

https://github.com/jsitnicki/kubecon-2024-sockmap/blob/main/examples/redir-bypass/redir_bypass.bpf.c

Section

Supported Socket Splicing Setups

Redirect

Redirect

send to local

Redirect use case → Bypass for containers

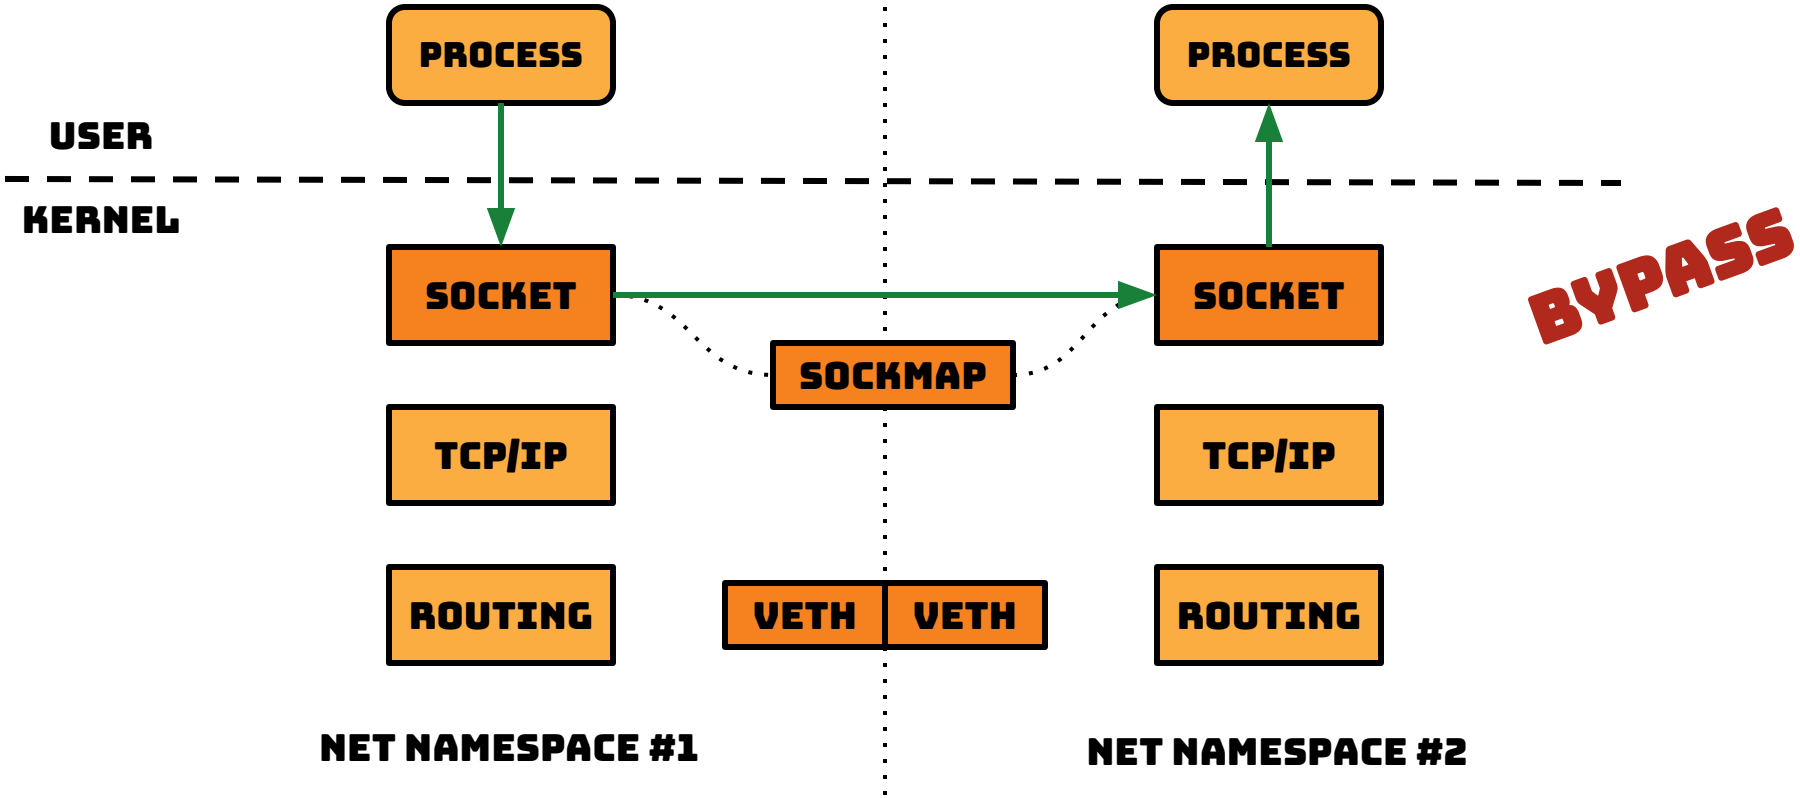

Redirect → send to local

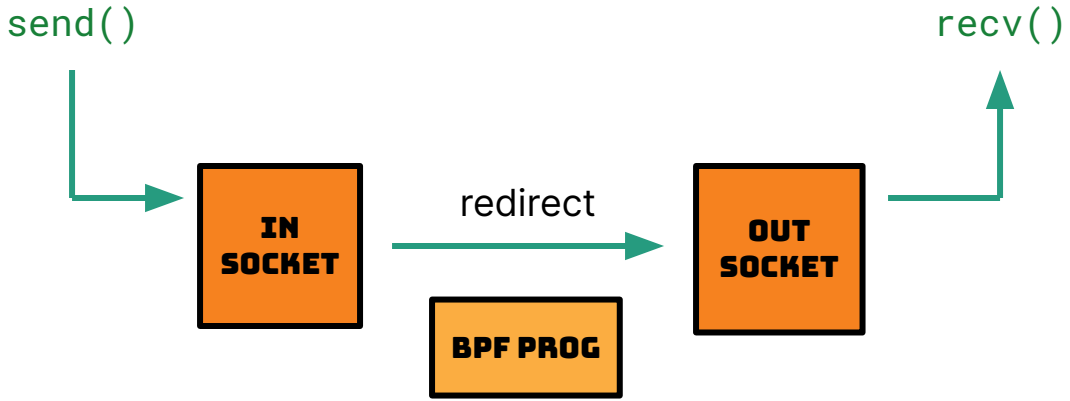

Like `socketpair()` or `pipe()`

Redirect → send to local → How?

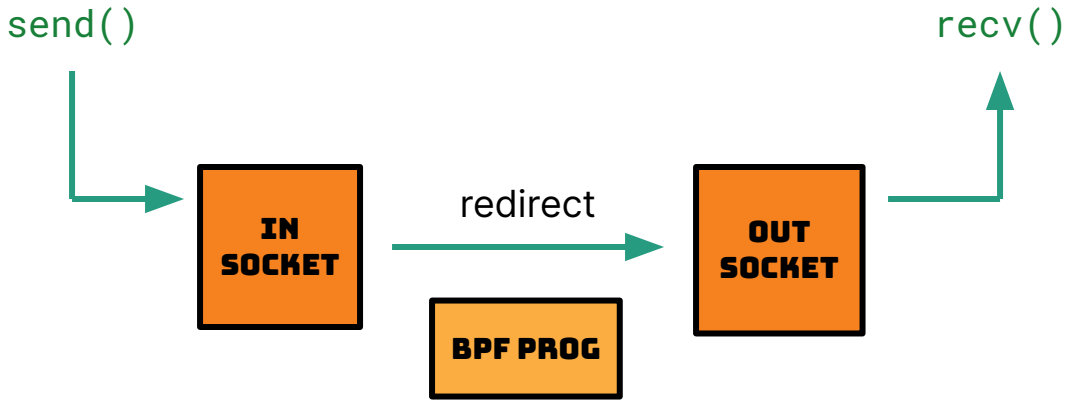

`BPF_PROG_TYPE_SK_MSG` program

→ attached to `BPF_SK_MSG_VERDICT` hook

→ calls `bpf_msg_redirect_hash/map()` with `BPF_F_INGRESS` flag

→ returns `SK_PASS`

selects target socket

Redirect → send to local → Example

```
SEC("sk_msg")
int sk_msg_redir_ingress(struct sk_msg_md *msg)
{
    __u32 key = 0;

    if (msg->remote_port == bpf_htonl(53))
        key = 1;

     return bpf_msg_redirect_map(msg, &sockmap, key, BPF_F_INGRESS);
}
```

Redirect → send to local → What?

IN → OUT	TCP	UDP	UNIX STR	UNIX DGR	VSOCK STR	VSOCK SEQ
TCP	●	●	●	●	●	●
UDP	●	●	●	●	●	●
UNIX STR	●	●	●	●	●	●
UNIX DGR	●	●	●	●	●	●
VSOCK STR	●	●	●	●	●	●
VSOCK SEQ	●	●	●	●	●	●

TCP to any but VSOCK

Redirect

send to egress

Redirect → send to egress

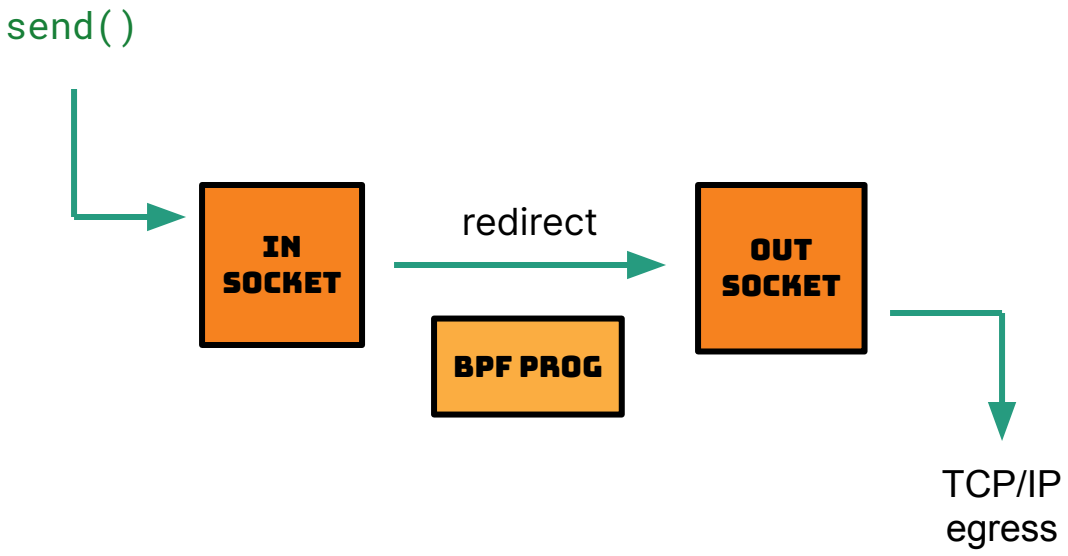

Sort of like `splice()`

(pipe → socket)

Redirect → send to egress → How?

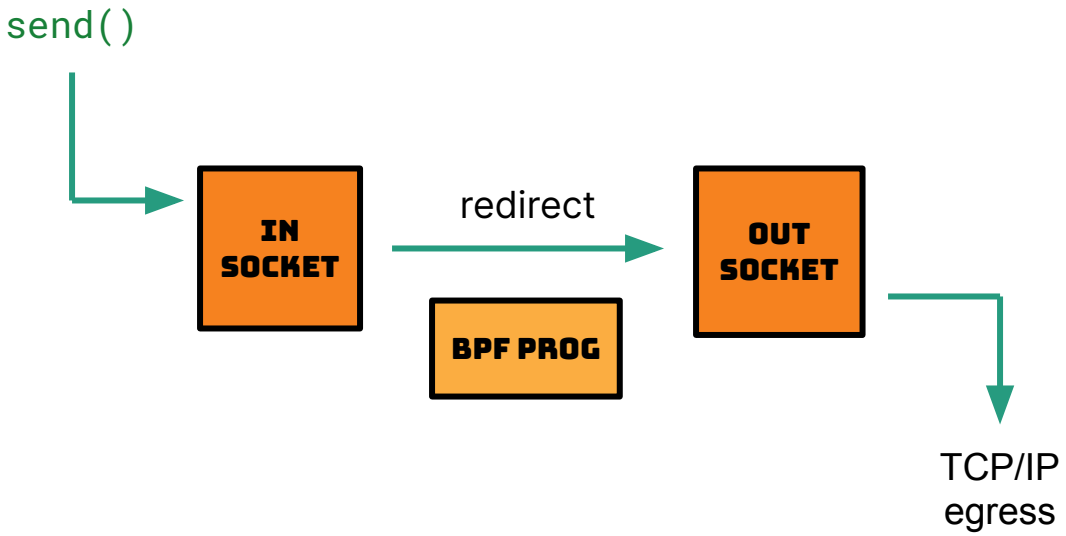

`BPF_PROG_TYPE_SK_MSG` prog

→ attached to `BPF_SK_MSG_VERDICT` hook

→ calls `bpf_msg_redirect_hash/map()` without any flags

→ returns `SK_PASS`

Redirect → send to egress → What?

IN → OUT	TCP	UDP	UNIX STR	UNIX DGR	VSOCK STR	VSOCK SEQ
TCP	●	●	●	●	●	●
UDP	●	●	●	●	●	●
UNIX STR	●	●	●	●	●	●
UNIX DGR	●	●	●	●	●	●
VSOCK STR	●	●	●	●	●	●
VSOCK SEQ	●	●	●	●	●	●

TCP to TCP only

Redirect

ingress to egress

Redirect use case → L7 network proxy

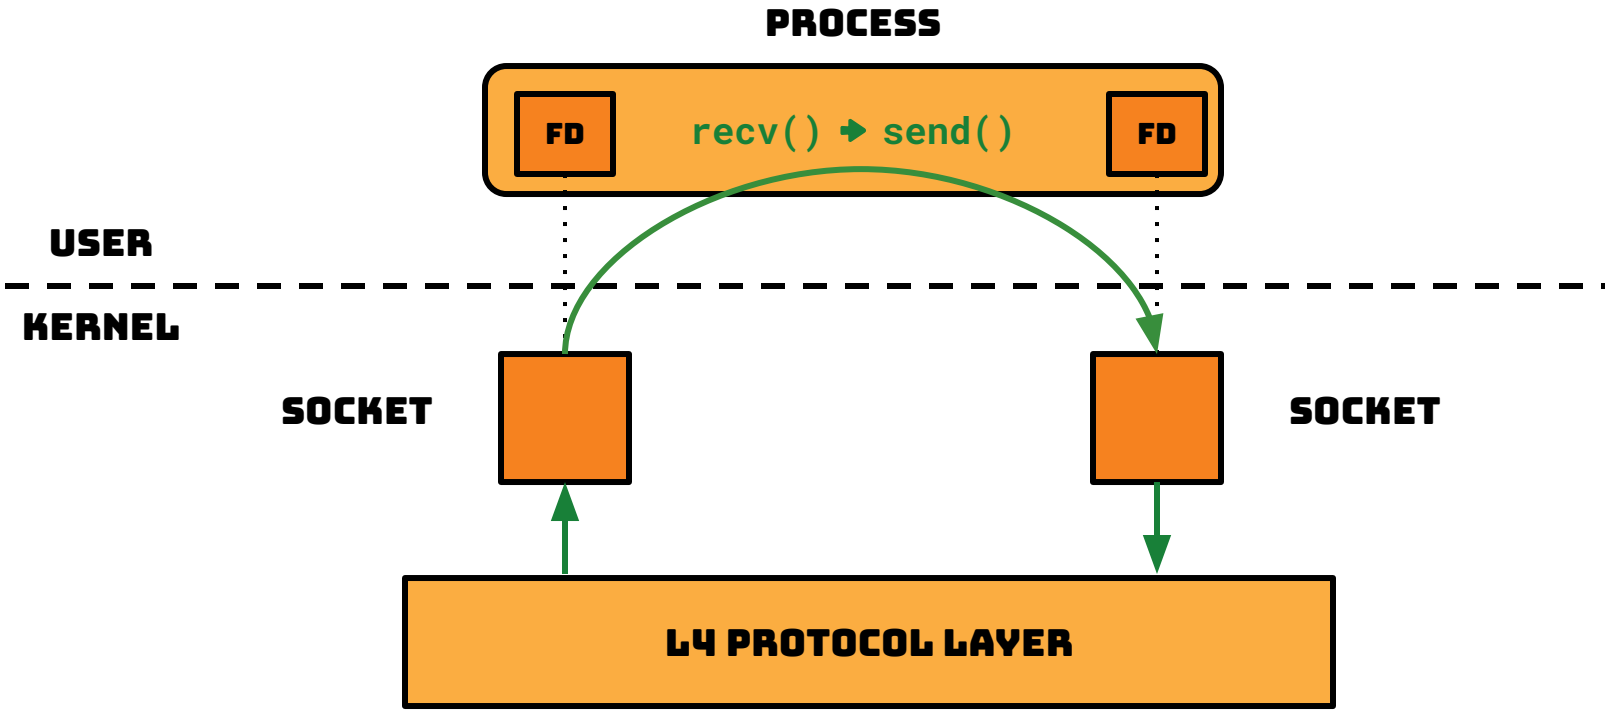

Examples: `socat`, `systemd-socket-proxyd`

Redirect use case → L7 network proxy

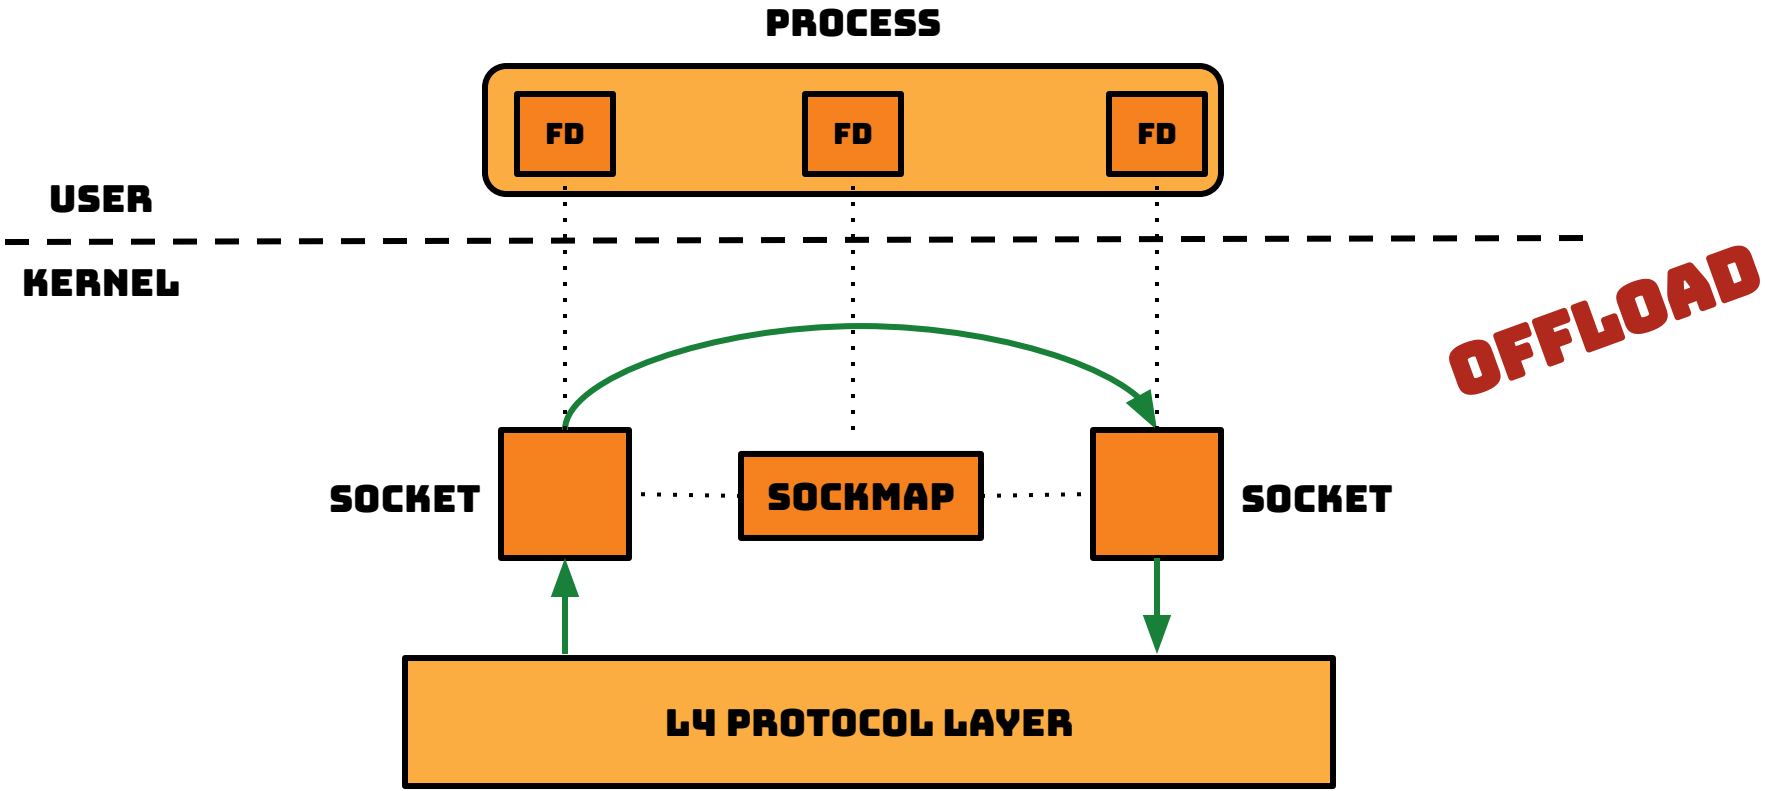

Redirect → ingress to egress

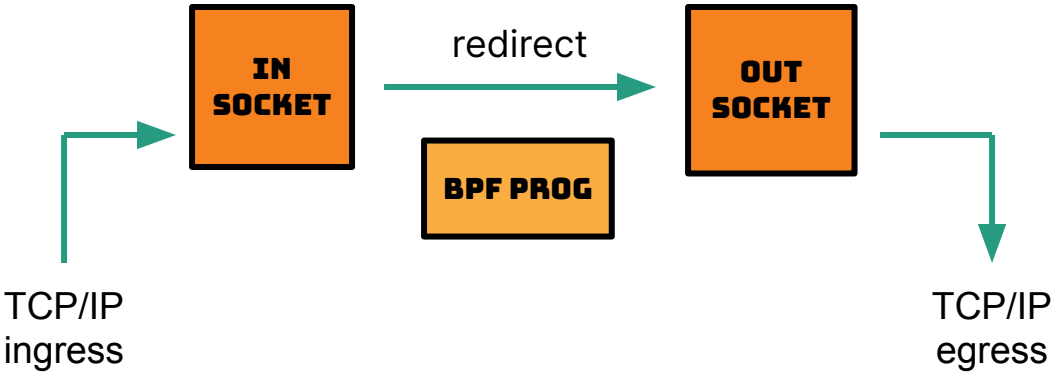

Like double `splice()`

(socket → pipe → socket)

Redirect → ingress to egress → How?

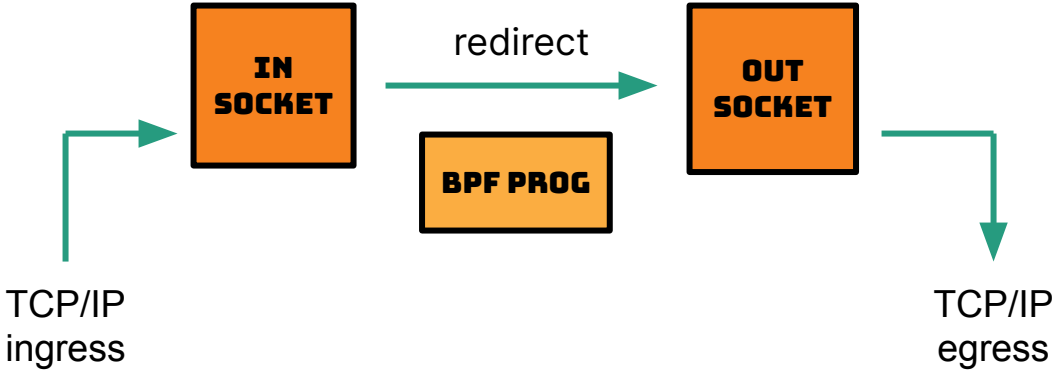

`BPF_PROG_TYPE_SK_SKB` prog

→ attached to `BPF_SK_SKB_VERDICT` hook

→ calls `bpf_sk_redirect_hash/map()` without any flags

→ returns `SK_PASS`

Redirect → ingress to egress → What?

IN → OUT	TCP	UDP	UNIX STR	UNIX DGR	VSOCK STR	VSOCK SEQ
TCP	●	●	●	●	●	●
UDP	●	●	●	●	●	●
UNIX STR	●	●	●	●	●	●
UNIX DGR	●	●	●	●	●	●
VSOCK STR	●	●	●	●	●	●
VSOCK SEQ	●	●	●	●	●	●

any to any

Redirect

ingress to local

Redirect → ingress to local

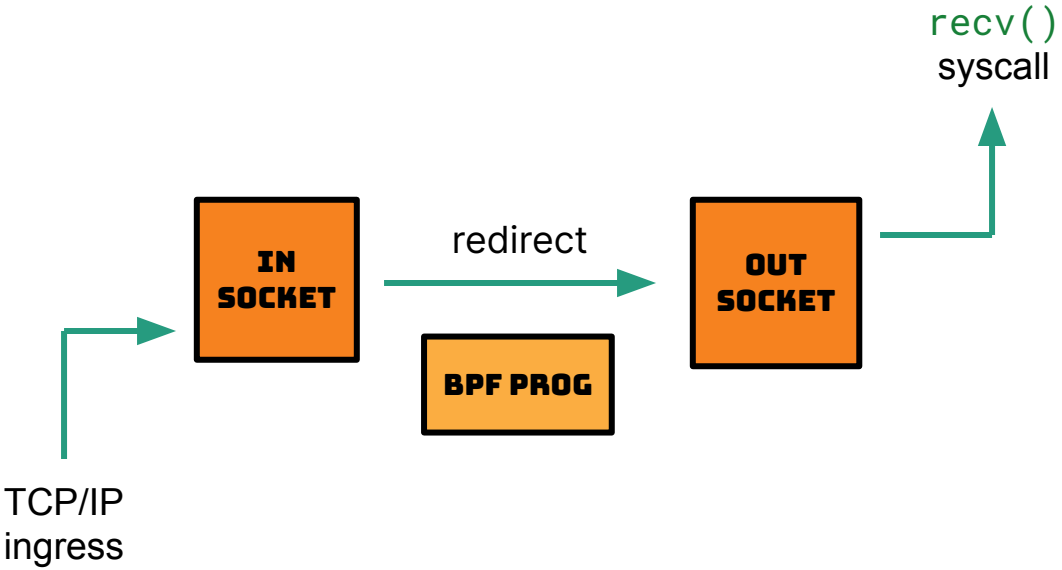

Sort of like `splice()`

(socket → pipe)

Redirect → ingress to local → How?

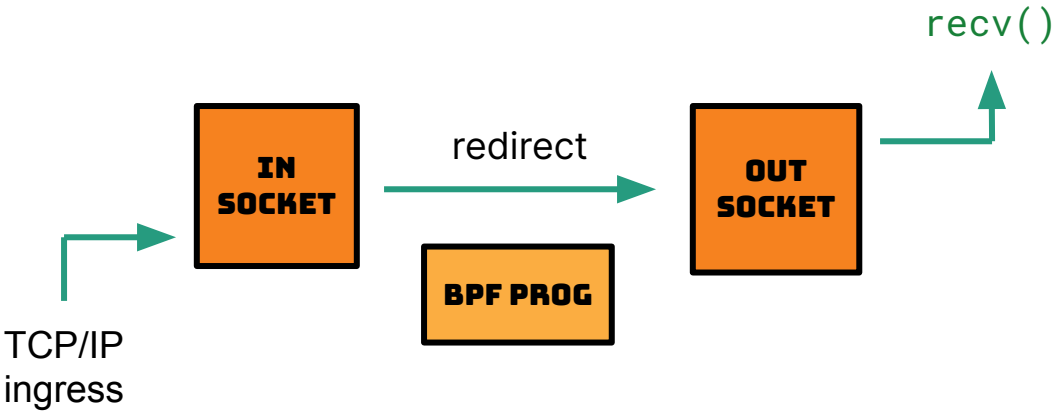

`BPF_PROG_TYPE_SK_SKB` prog

→ attached to `BPF_SK_SKB_VERDICT` hook

→ calls `bpf_sk_redirect_hash/map()` with `BPF_F_INGRESS` flag

→ returns `SK_PASS`

Redirect → ingress to local → What?

IN → OUT	TCP	UDP	UNIX STR	UNIX DGR	VSOCK STR	VSOCK SEQ
TCP	●	●	●	●	●	●
UDP	●	●	●	●	●	●
UNIX STR	●	●	●	●	●	●
UNIX DGR	●	●	●	●	●	●
VSOCK STR	●	●	●	●	●	●
VSOCK SEQ	●	●	●	●	●	●

any to any but VSOCK

Cheatsheet - Redirect with SOCKMAP

CLOUDFLARE

redirect scenario	program type BPF_PROG_TYPE_*	attach type BPF_*	redirect function	redirect flag	in socket type	out socket type
send to local	SK_MSG	SK_MSG_VERDICT	bpf_msg_redirect_*()	BPF_F_INGRESS	TCP	any but VSOCK
send to egress	SK_MSG	SK_MSG_VERDICT	bpf_msg_redirect_*()	none	TCP	TCP
ingress to egress	SK_SKB	SK_SKB_VERDICT	bpf_sk_redirect_*()	none	any	any
ingress to local	SK_SKB	SK_SKB_VERDICT	bpf_sk_redirect_*()	BPF_F_INGRESS	any	any but VSOCK

Section

Real life use-cases

Cilium project (CNI for K8S)

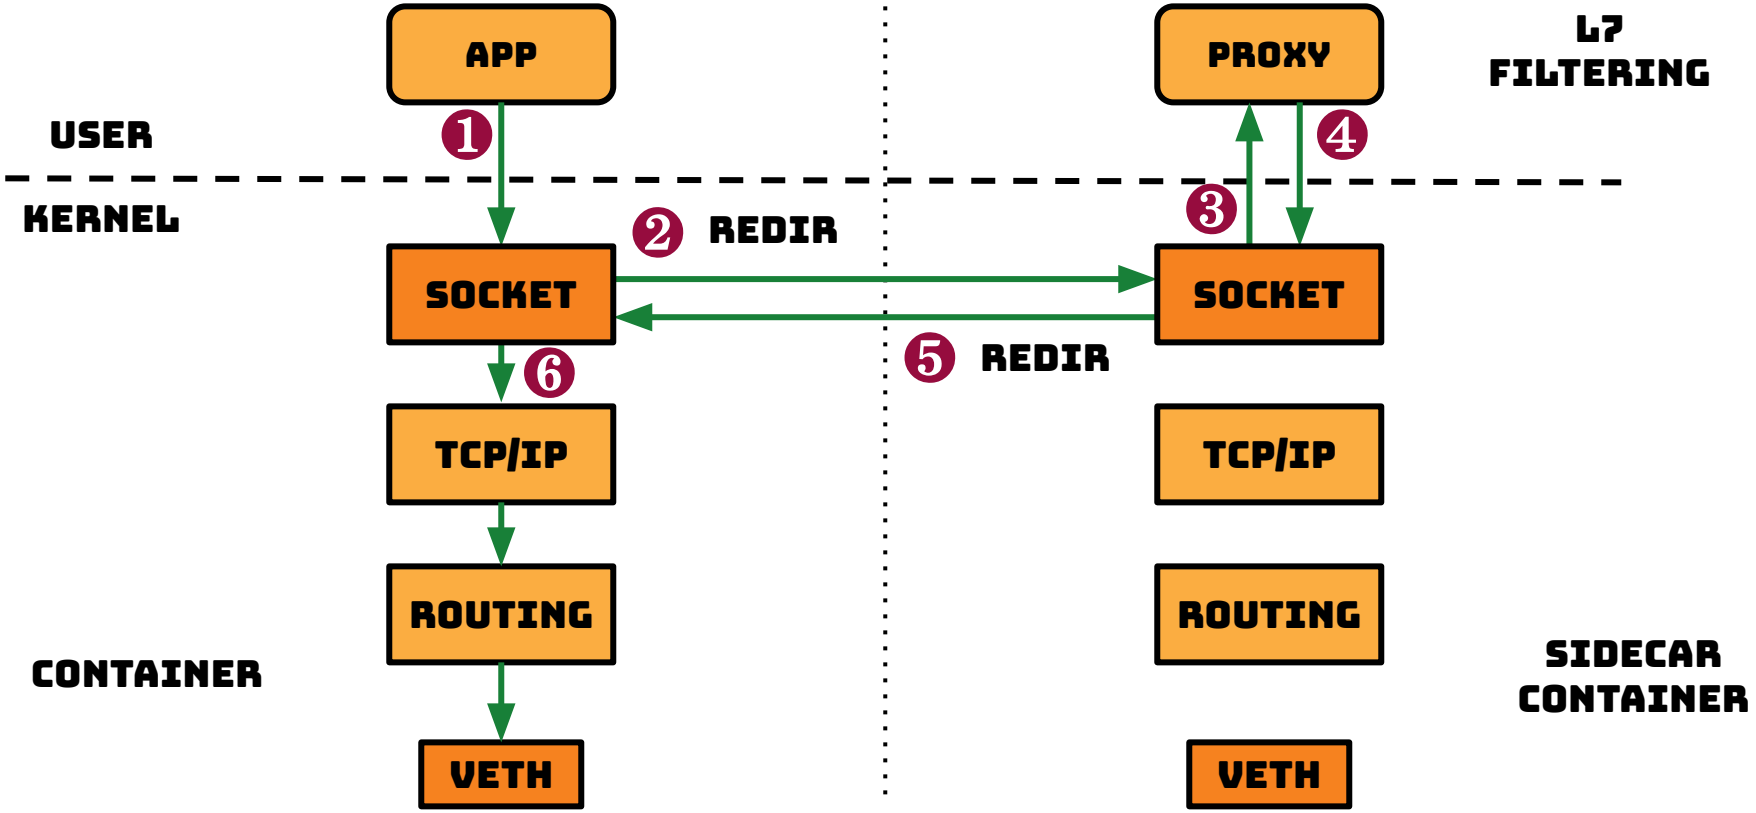

https://cilium.io/blog/2019/02/12/cilium-14/

Cilium project (CNI for K8S)

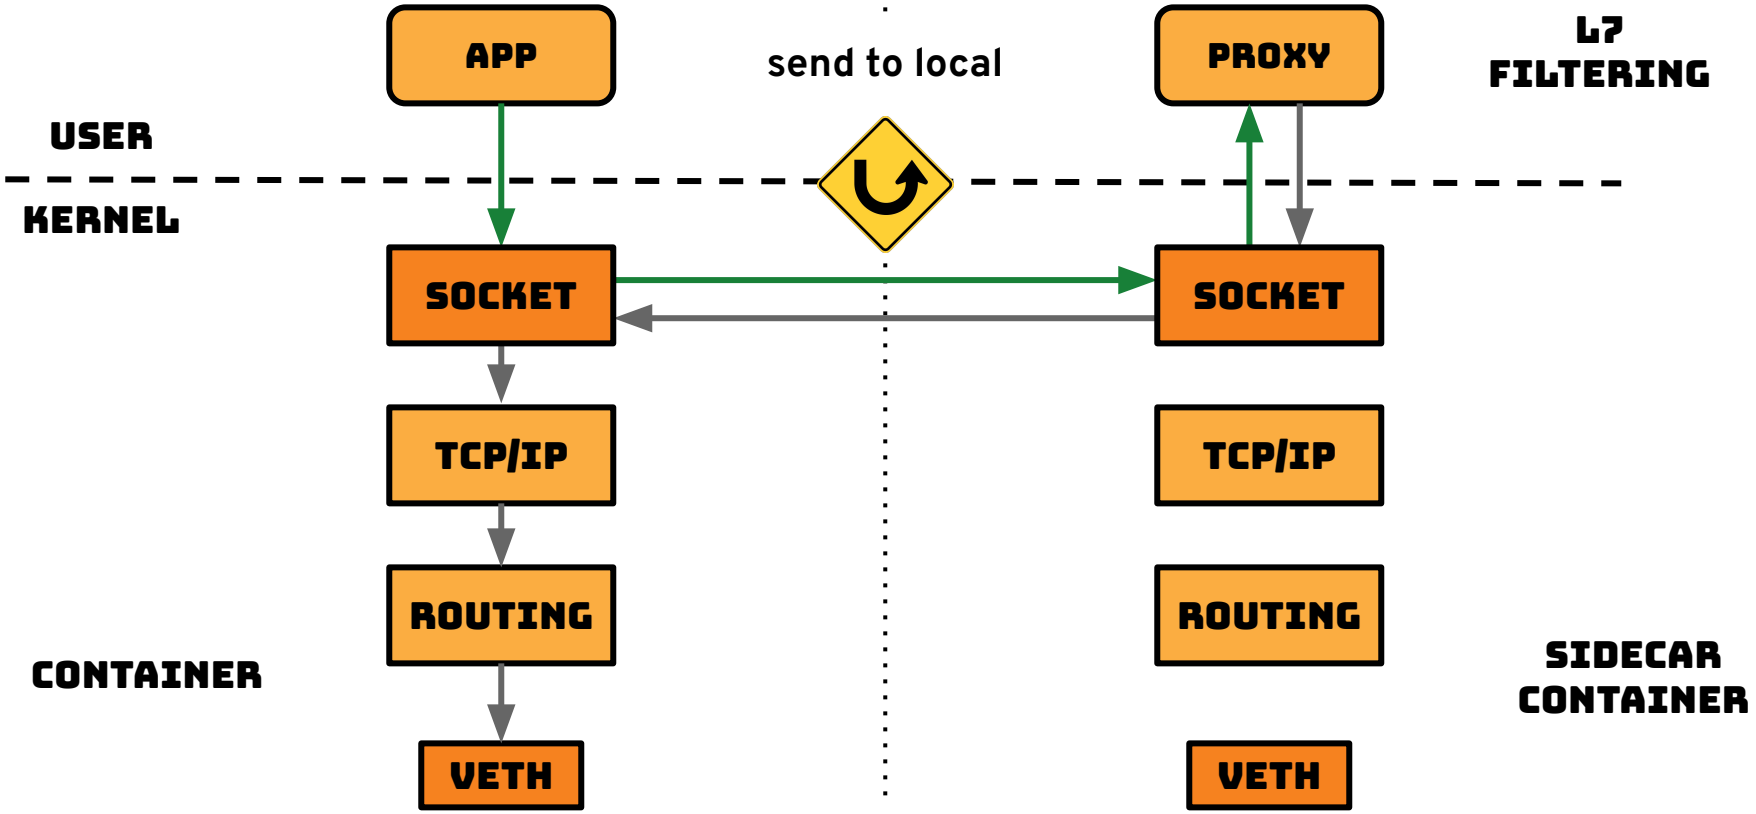

https://cilium.io/blog/2019/02/12/cilium-14/

Cilium project (CNI for K8S)

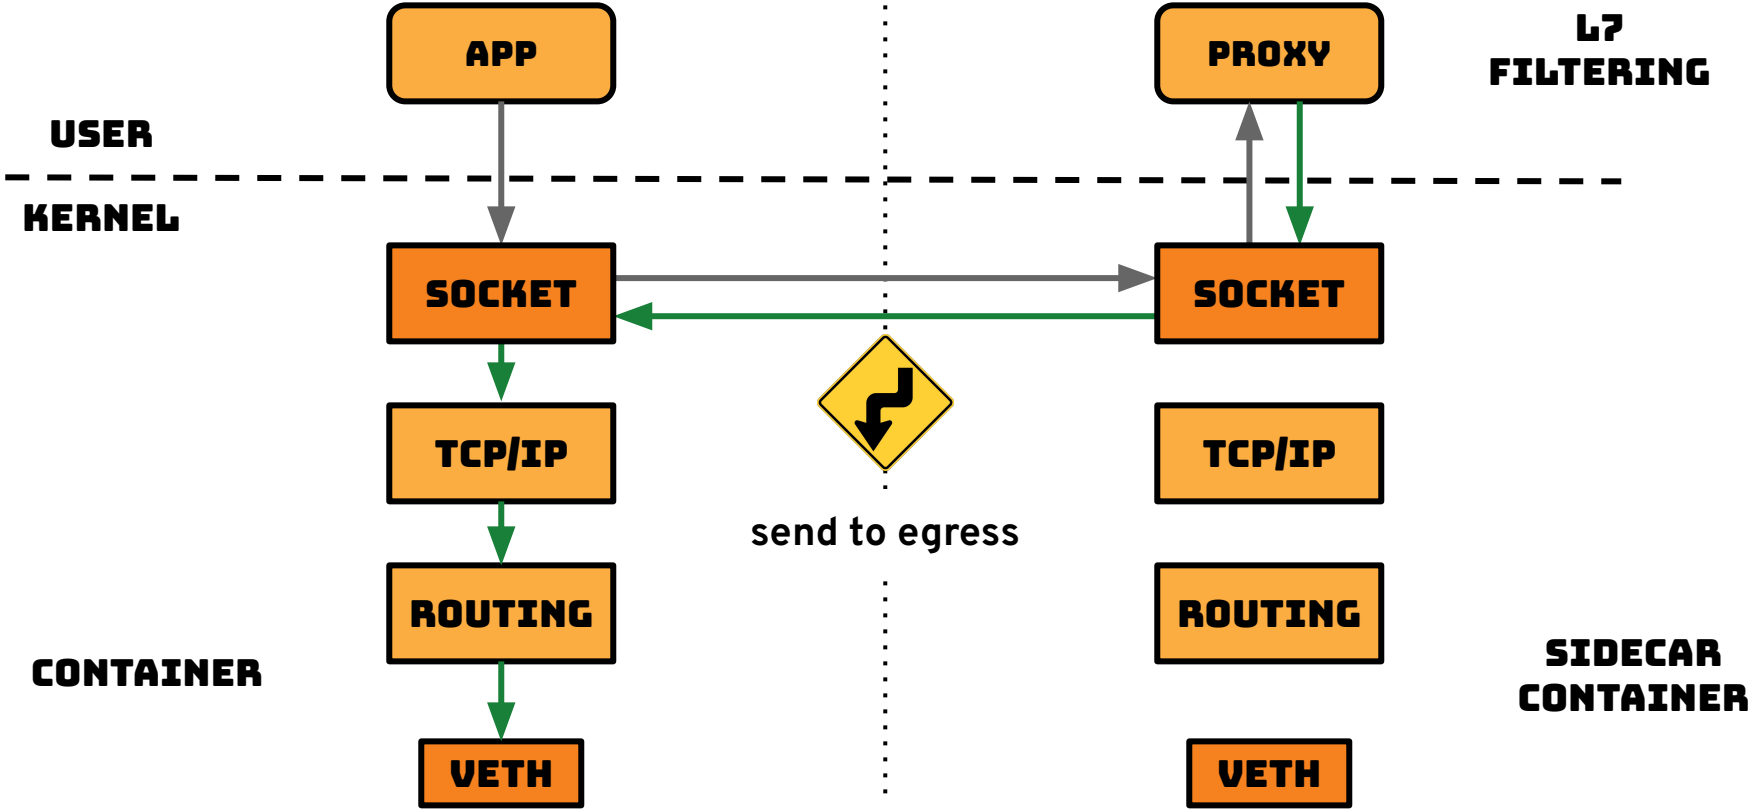

https://cilium.io/blog/2019/02/12/cilium-14/

Bytedance (TikTok)

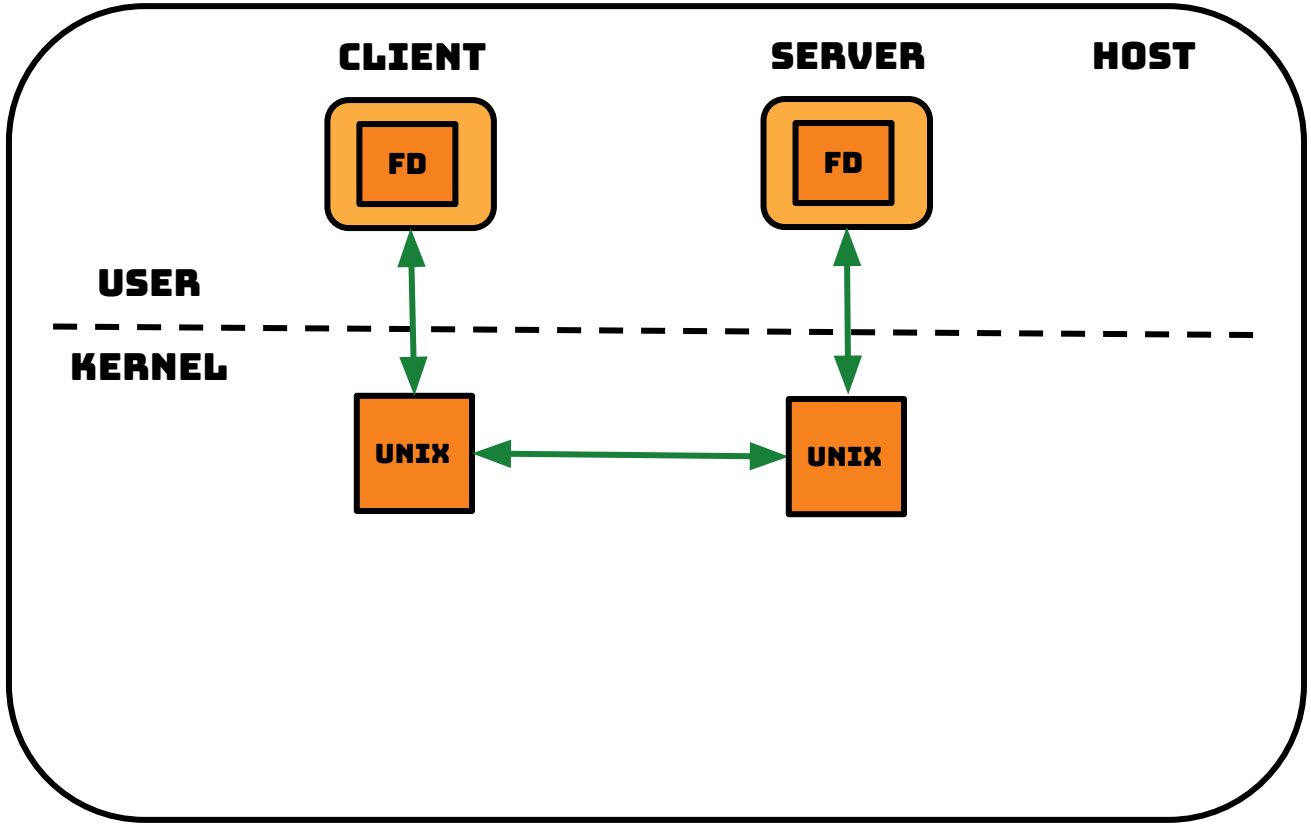

https://lore.kernel.org/bpf/20210331023237.41094-1-xiyou.wangcong@gmail.com/

Bytedance (TikTok)

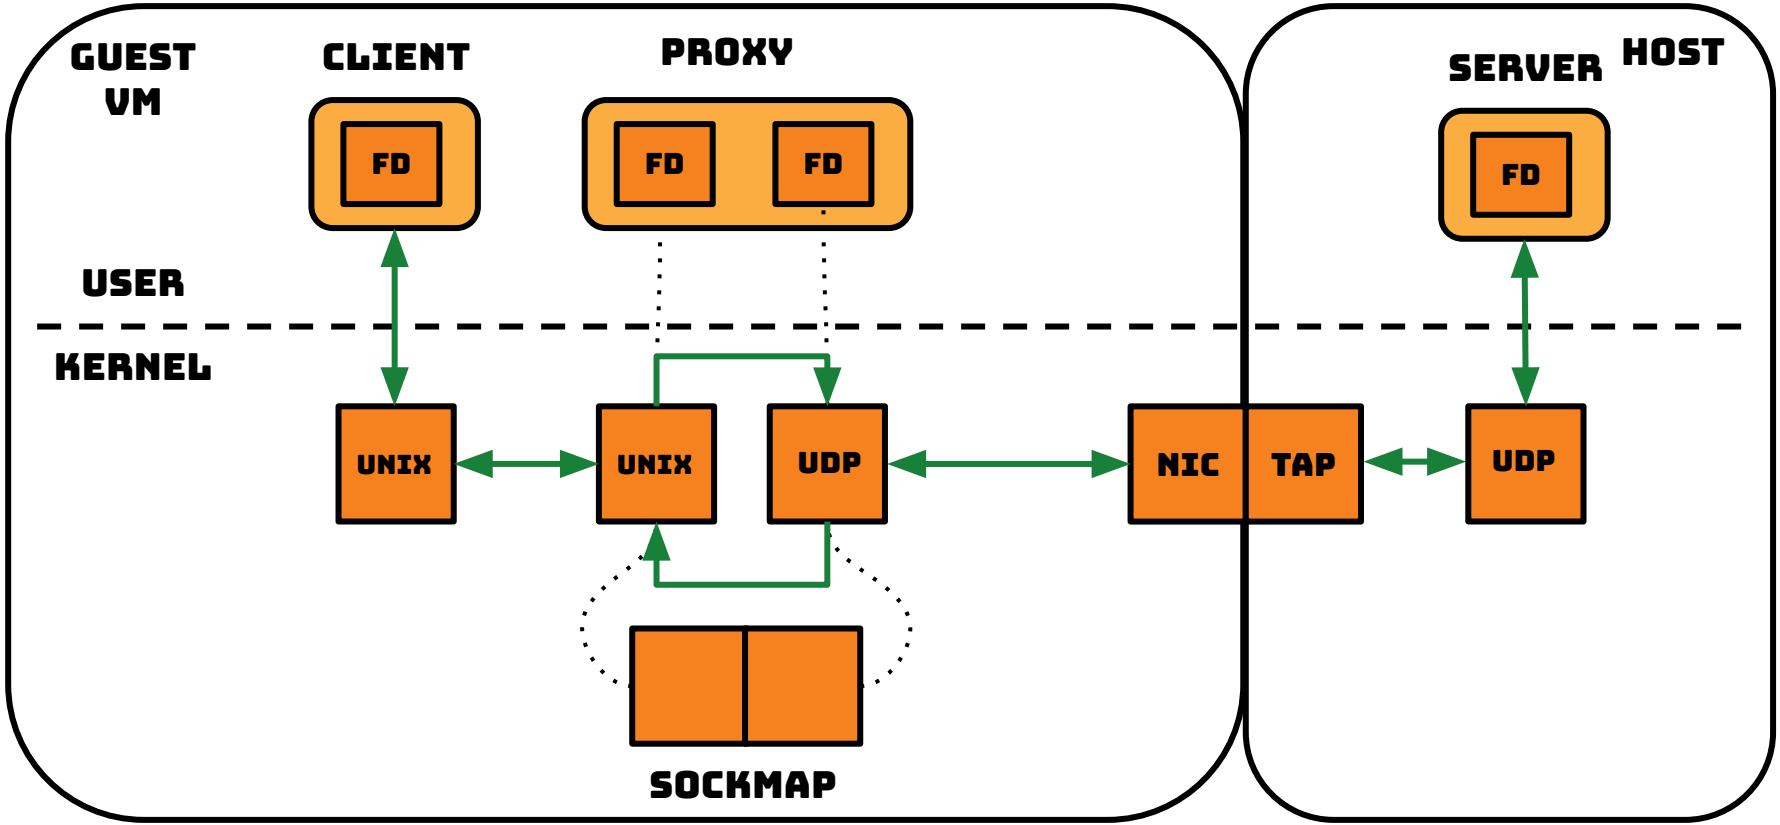

https://lore.kernel.org/bpf/20210331023237.41094-1-xiyou.wangcong@gmail.com/

Bytedance (TikTok) → Improved

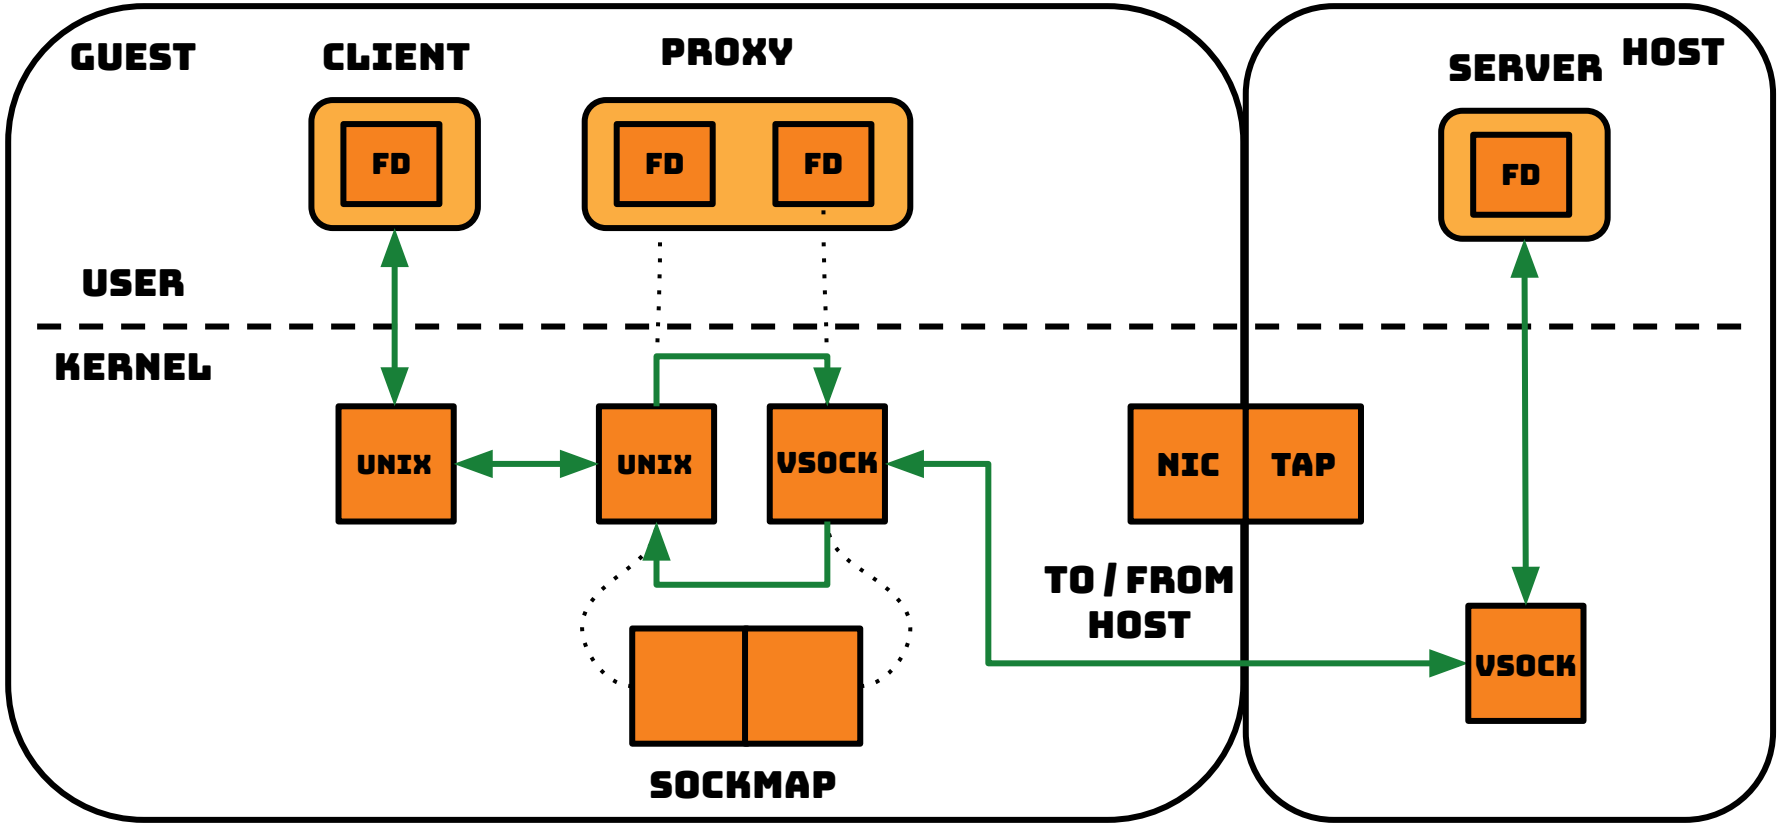

https://lore.kernel.org/all/20230327-vsock-sockmap-v4-0-c62b7cd92a85@bytedance.com/

Bytedance (TikTok) → Improved v2

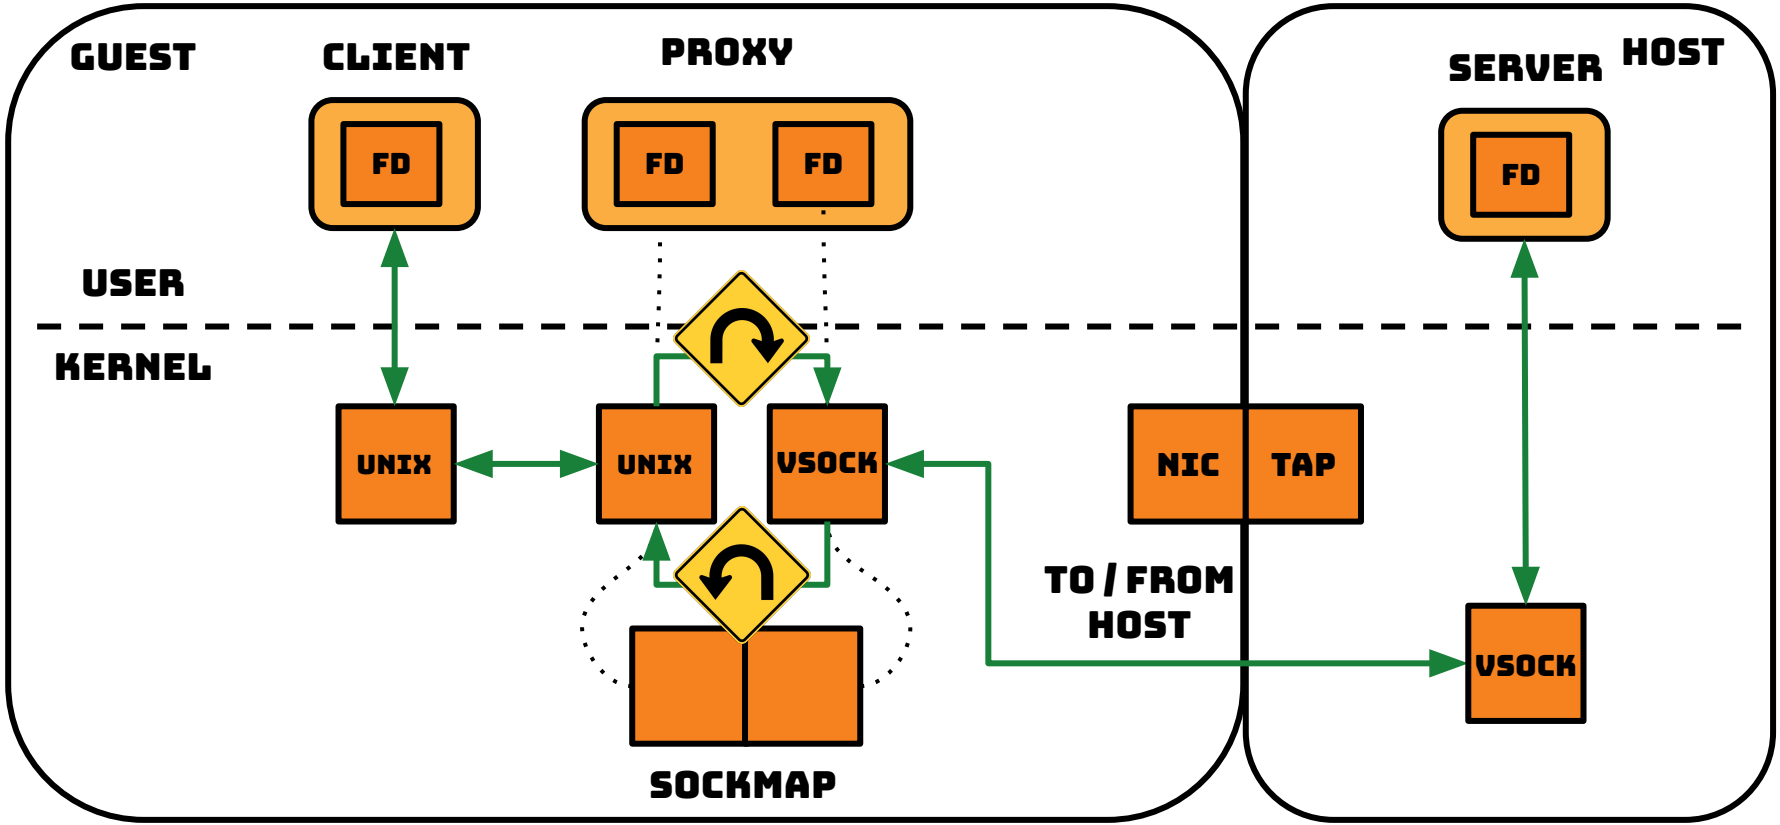

https://lore.kernel.org/all/20230327-vsock-sockmap-v4-0-c62b7cd92a85@bytedance.com/

Section

Where to learn more?

RESOURCES

CLOUDFLARE

1) **Linux Kernel → BPF Documentation → SOCKMAP and SOCKHASH map**

   Includes links to unit tests with API usage examples

2) **LPC 2018: Combining kTLS and BPF for Introspection and Policy Enforcement**

   See Daniel & John talk about Cilium SOCKMAP + kTLS use case (video, slides, paper)

3) **Cloudflare Blog: SOCKMAP - TCP splicing of the future**

   Read Marek review SOCKMAP from L7 proxy perspective

4) **eBPF Summit 2020: Steering connections to sockets with BPF socket lookup hook**

   Another use case for SOCKMAP as a container (video, slides, code)

THANK YOU

CLOUDFLARE

Reach out:

jakub@cloudflare.com

@jkbs0 @ ~~Twitter~~ X

jsitnicki @ LinkedIn

Mailing lists:

bpf@vger.kernel.org

netdev@vger.kernel.org

Slack:

#ebpf-kernel-dev @ cilium.slack.com

↑ code & slides repo ↑

https://github.com/jsitnicki/kubecon-2024-sockmap

Attribution

- Uncle Sam Wants You, Public Domain Pictures, CC0
- A bunch of question marks, Wikimedia Commons, CC BY-SA 4.0 DEED
- Nf knots, getarchive.net, CC0
- Construction blocks, vectorportal.com, CC BY 4.0 DEED
- Container crane image, vectorportal.com, CC BY 4.0 DEED
- Golden Gate bridge silhouette, rawpixel.com, CC0
- Evolution-des-wissens.jpg, Wikimedia Commons, CC BY-SA 3.0 DEED
- Leather pouch, rawpixel.com, CC0
- Redirection.svg, Wikimedia Commons, CC BY 3.0 DEED

Overflow slides

Section

EVOLUTION OF SOCKMAP

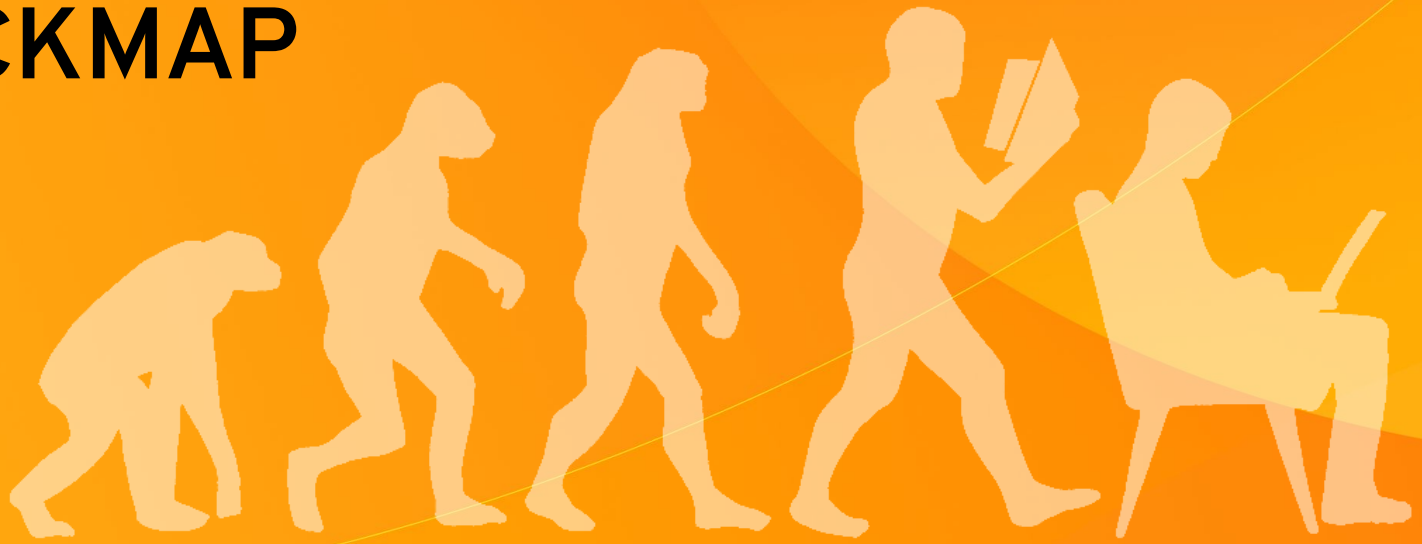

How did we get here?

Kernel Connection Multiplexor

Added infrastructure later reused by SOCKMAP - Psock and stream parser program.

2016
v4.6

```
+------------+   +------------+   +------------+   +------------+
| KCM socket |   | KCM socket |   | KCM socket |   | KCM socket |
+------------+   +------------+   +------------+   +------------+
      |                |               |                |
      +------------+   |               |     +----------+
                   |   |               |     |
               +-----------------------------------+
               |           Multiplexor             |
               +-----------------------------------+
                 |   |           |           |  |
       +---------+   |           |           |  ------------+
       |             |           |           |              |
+----------+  +----------+  +----------+  +----------+ +----------+
|  Psock   |  |  Psock   |  |  Psock   |  |  Psock   | |  Psock   |
+----------+  +----------+  +----------+  +----------+ +----------+
      |              |           |            |             |
+----------+  +----------+  +----------+  +----------+ +----------+
| TCP sock |  | TCP sock |  | TCP sock |  | TCP sock | | TCP sock |
+----------+  +----------+  +----------+  +----------+ +----------+
```

[PATCH v2 net-next 00/13] kcm: Kernel Connection Multiplexor (KCM)

How did we get here?

CLOUDFLARE

Kernel Connection Multiplexor

Added infrastructure later reused by SOCKMAP - Psock and stream parser program.

2016
v4.6

2017
v4.14

SOCKMAP initial version

Filtering and redirect on ingress to socket layer.

Counterpart of XDP DEVMAP.

author	John Fastabend <john.fastabend@gmail.com>	2017-08-15 22:32:47 -0700
committer	David S. Miller <davem@davemloft.net>	2017-08-16 11:27:53 -0700
commit	174a79ff9515f400b9a6115643dafd62a635b7e6 (patch)	
tree	f48f1fc407adb9bce6fb0e5cddaabd7141acd071	
parent	a6f6df69c48b86cd84f36c70593eb4968fceb34a (diff)	
download	linux-174a79ff9515f400b9a6115643dafd62a635b7e6.tar.gz	

bpf: sockmap with sk redirect support

```
Recently we added a new map type called dev map used to forward XDP
packets between ports (6093ec2dc313). This patches introduces a
similar notion for sockets.

A sockmap allows users to add participating sockets to a map. When
sockets are added to the map enough context is stored with the
map entry to use the entry with a new helper

  bpf_sk_redirect_map(map, key, flags)

This helper (analogous to bpf_redirect_map in XDP) is given the map
and an entry in the map. When called from a sockmap program, discussed
below, the skb will be sent on the socket using skb_send_sock().
```

commit 174a79ff9515 ("bpf: sockmap with sk redirect support")

How did we get here?

CLOUDFLARE

Kernel Connection Multiplexor

Added infrastructure later reused by SOCKMAP - Psock and stream parser program.

2016 v4.6

2017 v4.14

2018 v4.17

SK_MSG program

Filter and redirect at sendmsg() time (TCP only).

SOCKMAP initial version

Filter and redirect on ingress to socket layer (TCP only).

Counterpart of XDP DEVMAP.

author	John Fastabend <john.fastabend@gmail.com>	2018-03-18 12:57:10 -0700
committer	Daniel Borkmann <daniel@iogearbox.net>	2018-03-19 21:14:38 +0100
commit	4f738adba30a7cfc006f605707e7aee847ffefa0 (patch)	
tree	6603749a44356d3a44110c44f890a45b88d7e935	
parent	8c05dbf04b2882c3c0bc43fe7668c720210877f3 (diff)	
download	linux-4f738adba30a7cfc006f605707e7aee847ffefa0.tar.gz	

bpf: create tcp_bpf_ulp allowing BPF to monitor socket TX/RX data

```
This implements a BPF ULP layer to allow policy enforcement and
monitoring at the socket layer. In order to support this a new
program type BPF_PROG_TYPE_SK_MSG is used to run the policy at
the sendmsg/sendpage hook. To attach the policy to sockets a
sockmap is used with a new program attach type BPF_SK_MSG_VERDICT.
```


commit 4f738adba30a ("bpf: create tcp_bpf_ulp allowing BPF to monitor socket TX/RX data")

How did we get here?

CLOUDFLARE

2016 v4.6

Kernel Connection Multiplexor

Added infrastructure later reused by SOCKMAP - Psock and stream parser program.

2017 v4.14

SOCKMAP initial version

Filter and redirect on ingress to socket layer (TCP only).

Counterpart of XDP DEVMAP.

2018 v4.17

SK_MSG program

Filter and redirect at sendmsg() time (TCP only).

2018 v4.18

SOCKHASH map

Same as SOCKMAP but with flexible lookup key - blob of bytes, e.g. 5 tuple.

Backed by a hash table.

How did we get here?

CLOUDFLARE

2016 v4.6

Kernel Connection Multiplexor

Added infrastructure later reused by SOCKMAP - Psock and stream parser program.

2017 v4.14

SOCKMAP initial version

Filter and redirect on ingress to socket layer (TCP only).

Counterpart of XDP DEVMAP.

2018 v4.17

SK_MSG program

Filter and redirect at sendmsg() time (TCP only).

2018 v4.18

SOCKHASH map

Same as SOCKMAP but with flexible lookup key - blob of bytes, e.g. 5 tuple.

Backed by a hash table.

2018 v4.20

Switch to sk_msg API

Kernel TLS and SK_MSG can gracefully coexist.

Allows introspection / policy enforcement before in-kernel encryption.

commit 4f738adba30a ("bpf: create tcp_bpf_ulp allowing BPF to monitor socket TX/RX data")

How did we get here?

CLOUDFLARE

Store TCP or UDP sockets

SOCKMAP becomes a generic BPF map for sockets.

It can hold both connected and listening TCP sockets, and any bound UDP socket.

2020
v5.7

How did we get here?

CLOUDFLARE

Store TCP or UDP sockets

SOCKMAP becomes a generic BPF map for sockets.

It can hold both connected and listening TCP sockets, and any bound UDP socket.

2020
v5.7

2020
v5.10

More BPF programs can update SOCKMAP

Sockets can be inserted into SOCMAP by a few selected types of BPF programs.

Initially only SOCK_OPS programs could do it.

How did we get here?

CLOUDFLARE

Store TCP or UDP sockets

SOCKMAP becomes a generic BPF map for sockets.

It can hold both connected and listening TCP sockets, and any bound UDP socket.

2020 v5.7

2020 v5.10

More BPF programs can update SOCKMAP

Sockets can be inserted into SOCMAP by a few selected types of BPF programs.

Initially only SOCK_OPS programs could do it.

2020 v5.10

BPF iterators support

Iterate over SOCKMAP from BPF context.

Allows copying socket references from one SOCKMAP to another.

How did we get here?

CLOUDFLARE

2020 v5.7

Store TCP or UDP sockets

SOCKMAP becomes a generic BPF map for sockets.

It can hold both connected and listening TCP sockets, and any bound UDP socket.

2020 v5.10

More BPF programs can update SOCKMAP

Sockets can be inserted into SOCMAP by a few selected types of BPF programs.

Initially only SOCK_OPS programs could do it.

2020 v5.10

BPF iterators support

Iterate over SOCKMAP from BPF context.

Allows copying socket references from one SOCKMAP to another.

2021 v5.13

Redirect for UDP

Use new SK_SKB_VERDICT BPF program to redirect packets between UDP sockets.

How did we get here?

CLOUDFLARE

2020 v5.7

Store TCP or UDP sockets

SOCKMAP becomes a generic BPF map for sockets.

It can hold both connected and listening TCP sockets, and any bound UDP socket.

2020 v5.10

More BPF programs can update SOCKMAP

Sockets can be inserted into SOCMAP by a few selected types of BPF programs.

Initially only SOCK_OPS programs could do it.

2020 v5.10

BPF iterators support

Iterate over SOCKMAP from BPF context.

Allows copying socket references from one SOCKMAP to another.

2021 v5.13

Redirect for UDP

Use new SK_SKB_VERDICT BPF program to redirect packets between UDP sockets.

2021 v5.15

UNIX domain sockets

Allow storing and redirecting packets from/to UNIX sockets (dgram and stream).

How did we get here?

CLOUDFLARE

VSOCK domain sockets

Redirecting from/to VSOCK sockets (stream and seqpacket).

2023
v6.4

How did we get here?

CLOUDFLARE

VSOCK domain sockets

Redirecting from/to VSOCK sockets (stream and seqpacket).

Your contribution here

...

2023
v6.4

202x
v6.x

202x
v6.x

BPF_F_PERMANENT flag

Permanent redirects.

Run verdict program only once and remember the steering decision.

(Egress support only.)